The *open*EHR Archetype Model

Archetype Definition Language

Version 2 (ADL2)

Editors:{T Beale, S Heard}[1]

Revision: 2.1.0

Pages: 109

Keywords: EHR, health records, modelling, constraints, software

1. Ocean Informatics Australia

© 2003-2006 The *open*EHR Foundation

The *open*EHR foundation

is an independent, non-profit community, facilitating the creation and sharing of health records by consumers and clinicians via open-source, standards-based implementations.

Founding Chairman	David Ingram, Professor of Health Informatics, CHIME, University College London
Founding Members	Dr P Schloeffel, Dr S Heard, Dr D Kalra, D Lloyd, T Beale

email: info@openEHR.org **web**: http://www.openEHR.org

Copyright Notice

Amendment Record

Issue	Details	Raiser	Completed
	R E L E A S E 1.0.1		
2.1.0	**CR-000203**: Release 1.0 explanatory text improvements. Improve Archetype slot explanation.	T Beale	25 Jun 2006
	CR-000208: Improve ADL grammar for assertion expressions.	T Beale	
	CR-000160: Duration constraints. Added ISO 8601 patterns for duration in cADL.	S Heard	
	CR-000213: Correct ADL grammar for date/times to be properly ISO8601 compliant. Include 'T' in cADL patterns and dADL and cADL Date/time, Time and Duration values.	T Beale	
	CR-000216: Allow mixture of W, D etc in ISO8601 Duration (deviation from standard)	S Heard	
	R E L E A S E 1.0		
2.0.0	**CR-000136**. Add validity rules to ADL document.	T Beale	18 Jan 2006
	CR-000153. Synchronise ADL and AOM attribute naming.	T Beale	
	CR-000154. Convert ADL archetypes to dADL documents.	T Beale	
	CR-000171. Add validity check for cardinality & occurrences	A Maldondo	
	R E L E A S E 0.96		
1.3.0	**CR-000141**. Allow point intervals in ADL. Updated atomic types part of cADL section and dADL grammar section.	S Heard	18 Jun 2005
	CR-000142. Update dADL grammar to support assumed values.	T Beale	
	CR-000143. Add partial date/time values to dADL syntax.	T Beale	
	CR-000149. Add URIs to dADL and remove query() syntax.	T Beale	
	CR-000156. Update documentation of container types.	T Beale	
	CR-000138. Archetype-level assertions.	T Beale	
	R E L E A S E 0.95		
1.2.1	**CR-000125**. C_QUANTITY example in ADL manual uses old dADL syntax.	T Beale	11 Feb 2005
	CR-000115. Correct "/[xxx]" path grammar error in ADL. Create new section describing ADL path syntax.	T Beale	
	CR-000127. Restructure archetype specifications. Remove clinical constraint types section of document.	T Beale	

Issue	Details	Raiser	Completed
1.2	**CR-000110**. Update ADL document and create AOM document. Added explanatory material; added domain type support; rewrote of most dADL sections. Added section on assumed values, "controlled" flag, nested container structures. Change language handling. Rewrote OWL section based on input from: - University of Manchester, UK, - University Seville, Spain. Various changes to assertions due to input from the DSTC. Detailed review from Clinical Information Project, Australia. **Remove UML models to "Archetype Object Model" document.** Detailed review from UCL. **CR-000103**. Redevelop archetype UML model, add new keywords: allow_archetype, include, exclude. CR-000104. Fix ordering bug when use_node used. Required parser rules for identifiers to make class and attribute identifiers distinct. Added grammars for all parts of ADL, as well as new UML diagrams.	T Beale A Rector R Qamar I Román Martínez A Goodchild Z Z Tun E Browne T Beale T Austin T Beale K Atalag	15 Nov 2004
	R E L E A S E 0.9		
1.1	**CR-000079**. Change interval syntax in ADL.	T Beale	24 Jan 2004
1.0	**CR-000077**. Add cADL date/time pattern constraints. **CR-000078**. Add predefined clinical types. Better explanation of cardinality, occurrences and existence.	S Heard, T Beale	14 Jan 2004
0.9.9	**CR-000073**. Allow lists of Reals and Integers in cADL. **CR-000075**. Add predefined clinical types library to ADL. Added cADL and dADL object models.	T Beale, S Heard	28 Dec 2003
0.9.8	**CR-000070**. Create Archetype System Description. Moved Archetype Identification Section to new Archetype System document. Copyright Assgined by Ocean Informatics P/L Australia to The *open*EHR Foundation.	T Beale, S Heard	29 Nov 2003
0.9.7	Added simple value list continuation (",..."). Changed path syntax so that trailing '/' required for object paths. Remove ranges with excluded limits. Added terms and term lists to dADL leaf types.	T Beale	01 Nov 2003
0.9.6	Additions during HL7 WGM Memphis Sept 2003	T Beale	09 Sep 2003
0.9.5	Added comparison to other formalisms. Renamed CDL to cADL and dDL to dADL. Changed path syntax to conform (nearly) to Xpath. Numerous small changes.	T Beale	03 Sep 2003
0.9	Rewritten with sections on cADL and dDL.	T Beale	28 July 2003
0.8.1	Added basic type constraints, re-arranged sections.	T Beale	15 July 2003
0.8	Initial Writing	T Beale	10 July 2003

Trademarks

"Microsoft" and ".Net" are registered trademarks of the Microsoft Corporation.

"Java" is a registered trademark of Sun Microsystems.

"Linux" is a registered trademark of Linus Torvalds.

Acknowledgements

Sebastian Garde, Central Qld University, Australia, for german translations.

Table of Contents

1   Introduction

1.1   Purpose

This document describes the design basis and syntax of the Archetype Definition Language (ADL). It is intended for software developers, technically-oriented domain specialists and subject matter experts (SMEs). Although ADL is primarily intended to be read and written by tools, it is quite readable by humans and ADL archetypes can be hand-edited using a normal text editor.

The intended audience includes:

- Standards bodies producing health informatics standards;
- Software development organisations using *open*EHR;
- Academic groups using *open*EHR;
- The open source healthcare community;
- Medical informaticians and clinicians interested in health information;
- Health data managers.

1.2   Related Documents

Prerequisite documents for reading this document include:

- The *open*EHR Architecture Overview

Related documents include:

- The *open*EHR Archetype Object Model (AOM)
- The *open*EHR Archetype Profile (oAP)

1.3   Nomenclature

In this document, the term 'attribute' denotes any stored property of a type defined in an object model, including primitive attributes and any kind of relationship such as an association or aggregation. XML 'attributes' are always referred to explicitly as 'XML attributes'.

1.4   Status

This document is under development, and is published as a proposal for input to standards processes and implementation works.

This document is available at http://svn.openehr.org/specification/TAGS/Release-1.0/publishing/architecture/am/adl2.pdf.

The latest version of this document can be found at http://svn.openehr.org/specification/TRUNK/publishing/architecture/am/adl2.pdf.

Blue text indicates sections under active development.

1.5   Peer review

Known omissions or questions are indicated in the text with a "to be determined" paragraph, as follows:

TBD_1:   (example To Be Determined paragraph)

Areas where more analysis or explanation is required are indicated with "to be continued" paragraphs like the following:

To Be Continued: more work required

Reviewers are encouraged to comment on and/or advise on these paragraphs as well as the main content. Please send requests for information to info@openEHR.org. Feedback should preferably be provided on the mailing list openehr-technical@openehr.org, or by private email.

1.6   Conformance

Conformance of a data or software artifact to an *open*EHR Reference Model specification is determined by a formal test of that artifact against the relevant *open*EHR Implementation Technology Specification(s) (ITSs), such as an IDL interface or an XML-schema. Since ITSs are formal, automated derivations from the Reference Model, ITS conformance indicates RM conformance.

2 Overview

2.1 What is ADL?

Archetype Definition Language (ADL) is a formal language for expressing archetypes, which are constraint-based models of domain content. The archetype concept is described by Beale [1], [2]. The *open*EHR Archetype Object Model [3] describes an object model equivalent of the ADL syntax. The *open*EHR archetype framework is described in terms of Archetype Definitions and Principles [6] and an Archetype System [7]. Other semantic formalisms which were considered in the course of archetype, and some which remain relevant are described in detailed in section 11 on page 103.

To describe constraints on data which are instances of some information model (e.g. expressed in UML), ADL archetypes use three syntaxes: dADL (a data expression syntax), cADL (a constraint expression syntax), and a version of first-order predicate logic (FOPL). It is most useful when very generic information models are used for describing the data in a system, for example, where the logical concepts PATIENT, DOCTOR and HOSPITAL might all be represented using a small number of classes such as PARTY and ADDRESS. In such cases, archetypes are used to constrain the *valid* structures of instances of these generic classes to represent the desired domain concepts. In this way future-proof information systems can be built - relatively simple information models and database schemas can be defined, and archetypes supply the semantic modelling, completely outside the software. ADL can thus be used to write archetypes for any domain where formal object model(s) exist which describe data instances.

When archetypes are used at runtime in particular contexts, they are *composed* into larger constraint structures, with local or specialist constraints added, via the use of *templates*. The formalism of templates is dADL. Archetypes can be *specialised* by creating an archetypes that reference existing archetypes as parents; such archetypes can only make certain changes while remaining compatible with the parent.

Another major function of archetypes is to connect information structures to formal terminologies. Archetypes are language-neutral, and can be authored in and translated into any language.

Finally, archetypes are completely path-addressable in a manner similar to XML data, using path expressions that are directly convertible to Xpath expressions.

2.1.1 Structure

An ADL archetype is syntactically a dADL document, containing a definition section expression in the cADL syntax. The structure of an ADL archetype is shown in FIGURE 1.

```
archetype_id = <"some.archetype.id">
adl_version = <"2.0">
concept = <[concept_code]>
original_language = <"en-uk">
translations = <
  ...
>

description = <
  ...
>

declaration = <
        aADL: declaration
              statements
>
definition = <
        cADL:formal
             constraint
             definition
>
invariant = <
        aADL: assertion
              statements
>

ontology = <
...
>
revision_history = <
...
>
```

optional sections

plug-in syntaxes
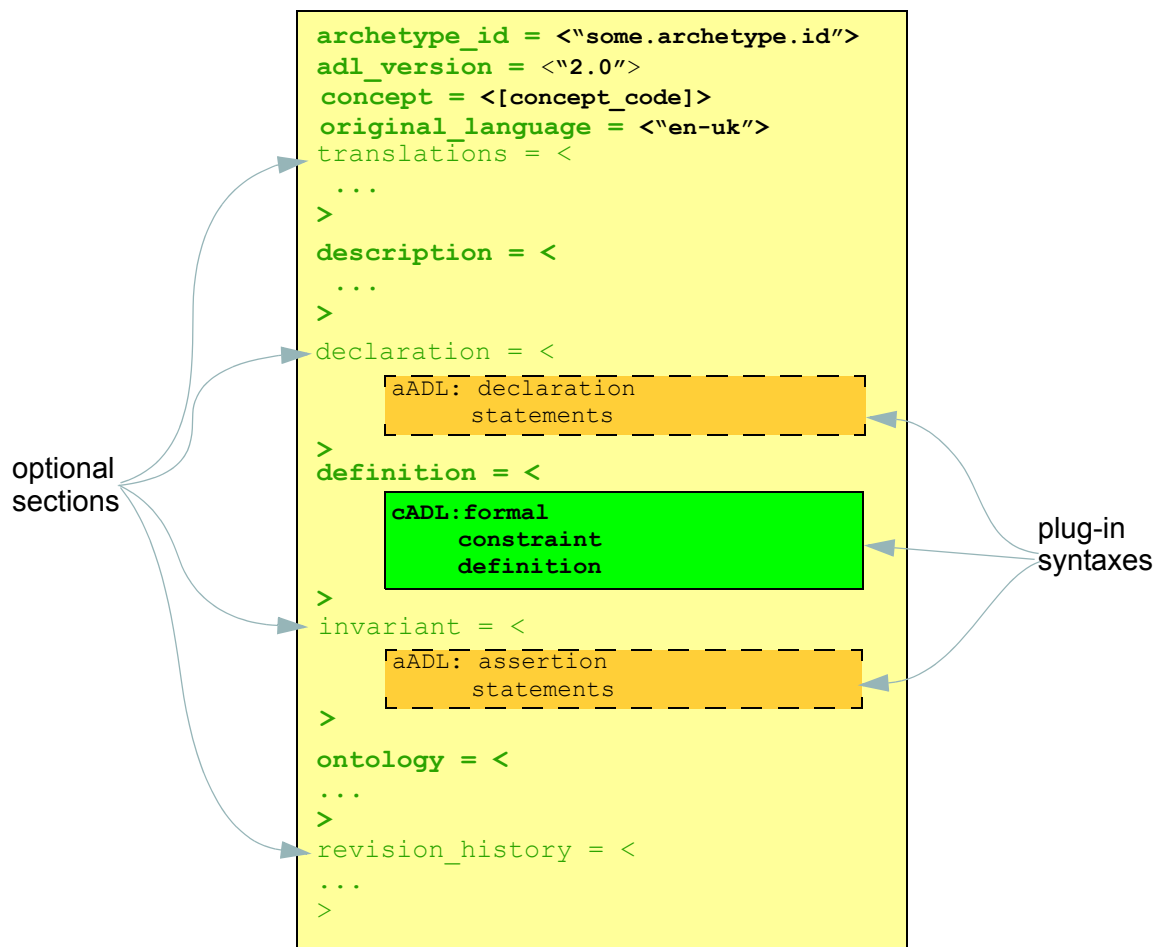
FIGURE 1 ADL Archetype Structure

This main part of this document describes dADL, cADL and ADL path syntax, before going on to describe the combined ADL syntax, archetypes and domain-specific type libraries.

2.1.2   An Example

The following is an example of a very simple archetype, giving a feel for the syntax. The main point to glean from the example is that the notion of 'guitar' is defined in terms of *constraints* on a *generic* model of the concept INSTRUMENT. The names mentioned down the left-hand side of the definition section ("INSTRUMENT", "size" etc) are alternately class and attribute names from an object model. Each block of braces encloses a specification for some particular set of instances that conform to a specific concept, such as 'guitar' or 'neck', defined in terms of constraints on types from a generic class model. The leaf pairs of braces enclose constraints on primitive types such as Integer, String, Boolean and so on.

```
archetype_id = <"adl-test-instrument.guitar.draft">
concept = <[at0000]> -- guitar
original_language = <"en">
definition = (cadl) <#
    INSTRUMENT[at0000] matches {
        size matches {|60..120|}                 -- size in cm
        date_of_manufacture matches {yyyy-mm-??}  -- year & month ok
```

```
            parts cardinality matches {0..*} matches {
                PART[at0001] matches {                    -- neck
                    material matches {[local::at0003,}  -- timber
                                             at0004]}   -- timber or nickel alloy
                }
                PART[at0002] matches {                    -- body
                    material matches {[local::at0003]}  -- timber
                }
            }
        }
    #>
    ontology = <
        term_definitions = <
            ["en"] = <
                ["at0000"] = <
                    text = <"guitar">;
                    description = <"stringed instrument">
                >
                ["at0001"] = <
                    text = <"neck">;
                    description = <"neck of guitar">
                >
                ["at0002"] = <
                    text = <"body">;
                    description = <"body of guitar">
                >
                ["at0003"] = <
                    text = <"timber">;
                    description = <"straight, seasoned timber">
                >
                ["at0004"] = <
                    text = <"nickel alloy">;
                    description = <"material for frets">
                >
            >
        >
    >
```

2.1.3   Semantics

As a parsable syntax, ADL has a formal relationship with structural models such as those expressed in UML, according to the scheme of FIGURE 2. Here we can see that ADL documents are parsed into a network of objects (often known as a 'parse tree') which are themselves defined by a formal, abstract object model (see The openEHR Archetype Object Model (AOM)). Such a model can in turn be re-expressed as any number of concrete syntaxes, such as in a programming language, XML-schema or OMG IDL.

While ADL syntax remains the primary abstract formalism for expressing archetypes, the AOM defines the semantics of an archetype, in particular relationships which must hold true between the parts of an archetype for it to be valid as a whole.

2.2   Computational Context

An archetypes is a structured model of domain content, such as "blood pressure". Archetypes sit between knowledge resources in a computing environment, such as terminologies and ontologies, and runtime data in production systems. Their primary purpose is to provide a reusable, interoperable way
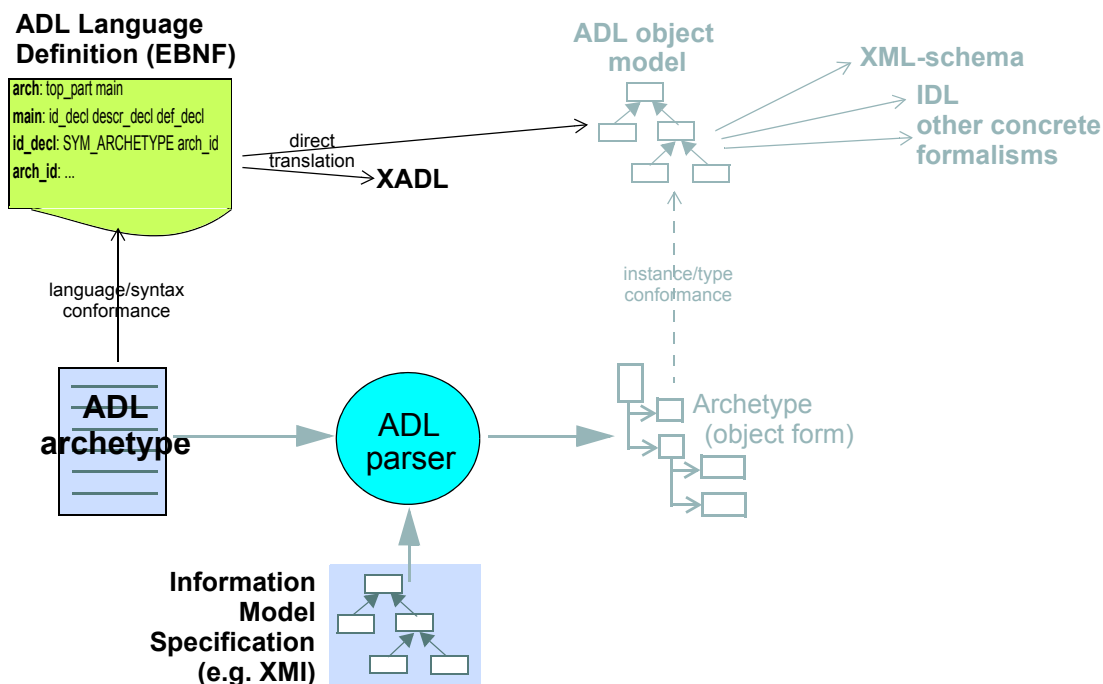
FIGURE 2 Relationship of ADL with Object Models

of managing data creation, validation and querying, by ensuring that data conform to particular struc-tures and semantic constraints. Every ADL archetype is written with respect to a particular informa-tion model, often known as a "reference model", if it is a shared, public specification.

Archetypes are applied to data via the use of *templates*, which are defined at a local level. Templates generally correspond closely to screen forms, and may be re-usable at a local or regional level. Tem-plates do not introduce any new semantics to archetypes, they simply specify the use of particular archetypes in particular hierarchical compositions, as well as default data values.

A third artifact that governs the functioning of archetypes and templates at runtime is a local *palette*, which specifies which natural language(s) and terminologies are in use in the locale. The use of a pal-ette removes irrelevant languages and terminology bindings from archetypes, retaining only those rel-evant to actual use. FIGURE 3 illustrates the overall environment in which archetypes, templates, and a locale palette exist.

2.3   XML form of Archetypes

With ADL parsing tools it is possible to convert ADL to any number of forms, including various XML formats. XML instance can be generated from the object form of an archetype in memory. An XML-schema corresponding to the ADL Object Model is published on *open*EHR.org.

2.4   Changes From Previous Versions

For existing users of ADL or archetype development tools, the following provides a guide to the changes in the syntax.

FIGURE 3 Archetypes, Templates, and Palettes

2.4.1 Version 2.0 from Version 1.3

Archetype is dADL Syntax

In this major revision of ADL, the ADL syntax is synchronised properly with the *open*EHR Archetype Object Model, although the semantics remain virtually unchanged. Accordingly, all top-level keywords in an ADL archetype are now considered the direct equivalent of an attribute of the same name in the AOM. To enable this, the entire ADL archetype is now a dADL document, with plug-in syntax sections for expressing constraints and invariants. The specific changes in this revision are:

- replace 'archetype' keyword with a dADL section named 'archetype_id' which has a String value;
- replace the ADL 1.2 adl_version syntax with a new top-level dADL section 'adl_version' which has a String value;
- replace the ADL 1.2 controlled syntax with a new top-level dADL section 'is_controlled' which has a Boolean value;
- replace 'specialise' keyword with a dADL section named 'parent_archetype_id';
- replace 'language' section with two top-level dADL sections 'original_language' and 'translations'.

These changes facilitate human understanding as well as automated processing, particularly serialisation/deserialisation into and out of object form.

A full dADL form of an archetype can also be supported, in which an archetype is a faithful dADL serialisation of instances of the Archetype Object Model (AOM), allowing archetypes to be parsed as dADL documents. This makes conversion to and from various XML formats trivial.

Archetype Slot syntax

In ADL2, two kinds of reference are allowed in archetype slots: archetype object references, e.g.

```
@archetype_id ∈ {/.*\.iso-ehr\.section\..*\..*/}
```

and archetype path references, e.g.

```
∃ /subject/relation
```

2.4.2   Version 1.3 from Version 1.2

The specific changes made in version 1.3 of ADL are as follows.

Query syntax replaced by URI data type

In version 1.2 of ADL, it was possible to include an external query, using syntax of the form:

```
attr_name = <query("some_service", "some_query_string")>
```

This is now replaced by the use of URIs, which can express queries, for example:

```
attr_name = <http://some.service.org?some%20query%20etc>
```

No assumption is made about the URI; it need not be in the form of a query - it may be any kind of URI.

Top-level Invariant Section

In this version, invariants can only be defined in a top level block, in a way similar to object-oriented class definitions, rather than on every block in the definition section, as is the case in version 1.2 of ADL. This simplifies ADL and the Archetype Object Model, and makes an archetype more comprehensible as a "type" definition.

2.4.3   Version 1.2 from Version 1.1

ADL Version

The ADL version is now optionally (for the moment) included in the first line of the archetype, as follows.

```
archetype (adl_version=1.2)
```

It is strongly recommended that all tool implementors include this information when archetypes are saved, enabling archetypes to gradually become imprinted with their correct version, for more reliable later processing. The adl_version indicator is likely to become mandatory in future versions of ADL.

dADL Syntax Changes

The dADL syntax for container attributes has been altered to allow paths and typing to be expressed more clearly, as part of enabling the use of Xpath-style paths. ADL 1.1 dADL had the following appearance:

```
school_schedule = <
    locations(1) = <...>
    locations(2) = <...>
    locations(3) = <...>
    subjects("philosophy:plato") = <...>
    subjects("philosophy:kant") = <...>
    subjects("art") = <...>
>
```

This has been changed to look like the following:

```
school_schedule = <
    locations = <
        [1] = <...>
        [2] = <...>
        [3] = <...>
    >
    subjects = <
        ["philosophy:plato"] = <...>
        ["philosophy:kant"] = <...>
```

```
    ["art"] = <...>
>
```

The new appearance both corresponds more directly to the actual object structure of container types, and has the property that paths can be constructed by directly reading identifiers down the backbone of any subtree in the structure. It also allows the optional addition of typing information anywhere in the structure, as shown in the following example:

```
school_schedule = SCHEDULE <
    locations = LOCATION <
        [1] = <...>
        [2] = <...>
        [3] = ARTS_PAVILLION <...>
    >
    subjects = <
        ["philosophy:plato"] = ELECTIVE_SUBJECT <...>
        ["philosophy:kant"] = ELECTIVE_SUBJECT <...>
        ["art"] = MANDATORY_SUBJECT <...>
    >
```

These changes will affect the parsing of container structures and keys in the description and ontology parts of the archetype.

Revision History Section

Revision history is now recorded in a separate section of the archetype, both to logically separate it from the archetype descriptive details, and to facilitate automatic processing by version control systems in which archtypes may be stored. This section is included at the end of the archetype because it is in general a monotonically growing section.

Primary_language and Languages_available Sections

An attribute previously called 'primary_language' was required in the ontology section of an ADL 1.1 archetype. This is renamed to 'original_language' and is now moved to a new top level section in the archetype called 'language'. Its value is still expressed as a dADL String attribute. The 'languages_available' attribute previously required in the ontology section of the archetype is renamed to 'translations', no longer includes the original languages, and is also moved to this new top level section.

2.5   Tools

A validating ADL parser is freely available from http://my.openehr.org/wsvn/oe_distrib/. It has been wrapped for use in Java and Microsoft .Net, and standard C/C++ environments. See the website for the latest status.

3   dADL - Data ADL

3.1   Overview

The dADL syntax provides a formal means of expressing *instance data* based on an underlying information model, which is readable both by humans and machines. The general appearance is exemplified by the following:

```
person = (List<PERSON>) <
    [01234] = <
        name = <                                  -- person's name
            forenames =    <"Sherlock">
            family_name =  <"Holmes">
            salutation =   <"Mr">
        >
        address = <                               -- person's address
            habitation_number = <"221B">
            street_name =  <"Baker St">
            city =         <"London">
            country =      <"England">
        >
    >
    [01235] = <
        -- etc
    >
>
```

In the above the attribute names `person`, `name`, `address` etc, and the type `List<PERSON>` are all assumed to come from an information model. The `[01234]` and `[01235]` tags identify container items.

The basic design principle of dADL is to be able to represent data in a way that is both machine-processable and human readable, while making the fewest assumptions possible about the information model to which the data conforms. To this end, type names are optional; often, only attribute names and values are explicitly shown. No syntactical assumptions are made about whether the underlying model is relational, object-oriented or what it actually looks like. More than one information model can be compatible with the same dADL-expressed data. The UML semantics of composition/aggregation and association are expressible, as are shared objects. Literal leaf values are only of 'standard' widely recognised types, i.e. Integer, Real, Boolean, String, Character and a range of Date/time types. In standard dADL, all complex types are expressed structurally.

A common question about dADL is why it is needed, when there is already XML? To start with, this question highlights the widespread misconception about XML, namely that because it can be read by a text editor, it is intended for humans. In fact, XML is designed for machine processing, and is textual to guarantee its interoperability. Realistic examples of XML (e.g. XML-schema instance, OWL-RDF ontologies) are generally unreadable for humans. dADL is on the other hand designed as a human-writable and readable formalism that is also machine processable; it may be thought of as an *abstract syntax for object-oriented data*. dADL also differs from XML by:

- providing a more comprehensive set of leaf data types, including intervals of numerics and date/time types, and lists of all primitive types;
- adhering to object-oriented semantics, particularly for container types, which XML schema languages generally do not;

- not using the confusing XML notion of 'attributes' and 'elements' to represent what are essentially object properties;
- requiring half the space of the equivalent XML.

Of course, this does not prevent XML exchange syntaxes being used for dADL, and indeed the conversion to XML instance is rather straighforward. Details on the XML expression of dADL and use of Xpath expressions is described in section 3.7 on page 35.

The dADL syntax as described above has a number of useful characteristics that enable the extensive use of paths to navigate it, and express references. These include:

- each <> block corresponds to an object (i.e. an instance of some type in an information model);
- the name before an '=' is always an attribute name or else a container element key, which attaches to the attribute of the enclosing block;
- paths can be formed by navigating down a tree branch and concatenating attribute name, container keys (where they are encountered) and '/' characters;
- every node is reachable by a path;
- shared objects can be referred to by path references.

3.2   Basics

3.2.1   Scope of a dADL Document
A dADL document may contain one or more objects from the same object model.

3.2.2   Keywords
dADL has no keywords of its own: all identifiers are assumed to come from an information model.

3.2.3   Reserved Characters
In dADL, a small number of characters are reserved and have the following meanings:

'<': open an object block;

'>': close an object block;

'=': indicate attribute value = object block;

'(', ')': type name or plug-in syntax type delimiters;

'<#': open an object block expressed in a plug-in syntax;

'#>': close an object block expressed in a plug-in syntax.

Within <> delimiters, various characters are used as follows to indicate primitive values:

'""': double quote characters are used to delimit string values;

'''': single quote characters are used to delimit single character values;

'|': bar characters are used to delimit intervals;

[]: brackets are used to delimit coded terms.

3.2.4 Comments

In a dADL text, comments satisfy the following rule:

comments are indicated by the characters "--". **Multi-line comments** are achieved using the "--" leader on each line where the comment continues. In this document, comments are shown in brown.

3.2.5 Reserved Character Encoding

Reserved characters need to be encoded in some contexts to avoid parsing ambiguity. Encoding is needed in strings, characters and URIs. The standard for encoding special characters, quotes, and other reserved characters is the Extensible Markup Language (XML) 1.0 (Third Edition) standard (http://www.w3.org/TR/REC-xml). Encoding of special characters according to this standard is as follows:

`&#NNN;` - decimal code for a character from ISO/IEC 10646

`&#xNNN;` - hexadecimal code for a character from ISO/IEC 10646

`&aaa;` - mnemonic code any character from ISO/IEC 10646, e.g. < means '<'

In dADL, special encoding is needed in the following situations:

in string values (e.g. "xxxx"): the '"' (double quote) character and possibly other special characters (defined in the XML standard) must be encoded;

in character values (e.g. 'x'): the ''' (single quote) character and possibly other special characters (defined in the XML standard) must be encoded;

in URIs: the <> characters, whitespace characters, and various others must be encoded according to the Uniform Resource Identifier (URI): Generic Syntax, Internet proposed standard RFC 3986, January 2005 (see http://www.ietf.org/rfc/rfc3986.txt).

In summary, special character encoding follows the rule:

Some reserved characters must be encoded within dADL String, Character and URI values, using the W3C XML standard or IETF RFC 3986 as appropriate.

3.2.6 Information Model Identifiers

Two types of identifiers from information models are used in dADL: type names and attribute names.

A **type name** is any identifier with an initial upper case letter, followed by any combination of letters, digits, and underscores. An **attribute name** is any identifier with an initial lower case letter, followed by any combination of letters, digits and underscores.

Type names in this document are in all uppercase, e.g. PERSON, except for 'built-in' types, such as primitive types (Integer, String, Boolean, Real, Double) and assumed container types (List<T>, Set<T>, Hash<T, U>), which are in mixed case, in order to provide easy differentiation of types assumed from constructed types.

Attribute names are shown in all lowercase, e.g. home_address. In both cases, underscores are used to represent word breaks. This convention is used to maximise the readability of this document, and other conventions may be used, such as the common programmer's mixed-case convention exemplified by Person and homeAddress, as long as they obey the rule above. The convention chosen for any particular dADL document should be based on the convention used in the underlying information model. Identifiers are shown in green in this document.

3.2.7 Semi-colons

Semi-colons can be used to separate dADL blocks, for example when it is preferable to include multiple attribute/value pairs on one line. Semi-colons make no semantic difference at all, and are included only as a matter of taste. The following examples are equivalent:

```
term = <text = <"plan">; description = <"The clinician's advice">>

term = <text = <"plan"> description = <"The clinician's advice">>

term = <
    text = <"plan">
    description = <"The clinician's advice">
>
```

Semi-colons are completely optional in dADL.

3.3 Paths

Because dADL data is hierarchical, and all nodes are uniquely identified, a reliable path can be determined for every node in a dADL text. The syntax of paths in dADL is the standard ADL path syntax, described in detail in section 7 on page 81. A path either finishes in a slash, and identifies an object node, or finishes in an attribute name, and identifies an attribute node. Paths are directly convertible to XPath expressions for use in XML-encoded data.

A typical ADL path used to refer to a node in a dADL text is as follows.

```
/term_definitions["en"]/items["at0001"]/text/
```

In the following sections, paths are shown for all the dADL data examples.

3.4 Structure

3.4.1 General Form

A dADL document expresses serialised instances of one or more complex objects. Each such instance is a hierarchy of attribute names and object values. In its simplest form, a dADL text consists of repetitions of the following pattern:

```
attribute_name = <value>
```

In the most basic form of dADL, each attribute name is the name of an attribute in an implied or actual object or relational model. Each "value" is either a literal value of a primitive type (see Primitive Types on page 31) or a further nesting of attribute names and values, terminating in leaf nodes of primitive type values. Where sibling attribute nodes occur, the attribute identifiers must be unique, just as in a standard object or relational model.

Sibling attribute names must be unique.

The following shows a typical structure.

```
attr_1 = <
    attr_2 = <
        attr_3 = <leaf_value>
        attr_4 = <leaf_value>
    >
    attr_5 = <
        attr_3 = <
```

```
            attr_6 = <leaf_value>
        >
        attr_7 = <leaf_value>
    >
>
attr_8 = <>
```

In the above structure, each "<>" encloses an instance of some type. The hierarchical structure corresponds to the part-of relationship between objects, otherwise known as *composition* and *aggregation* relationships in object-oriented formalisms such as UML (the difference between the two is usually described as being "sub-objects related by aggregation can exist on their own, whereas sub-objects related by composition are always destroyed with the parent"; dADL does not differentiate between the two, since it is the business of a model, not the data, to express such semantics). Associations between instances in dADL are also representable by references, and are described in section 3.4.6 on page 29.

3.4.1.1   Outer Delimiters

To be completely regular, an outer level of delimiters should be used, because the totality of a dADL text is an object, not a collection of disembodied attribute/object pairs. However, the outermost delimiters can be left out in order to improve readability, and without complicating the parsing process. The completely regular form would appear as follows:

```
<
    attr_1 = <
    >
    attr_8 = <>
>
```

   Outer '<>' delimiters in a dADL text are optional.

3.4.1.2   Paths

The complete set of paths for the above example is as follows.

```
attr_1
attr_1/attr_2
attr_1/attr_2/attr_3/        -- path to a leaf value
attr_1/attr_2/attr_4/        -- path to a leaf value
attr_1/attr_5
attr_1/attr_5/attr_3/
attr_1/attr_5/attr_3/attr_6  -- path to a leaf value
attr_1/attr_5/attr_7/        -- path to a leaf value
attr_8
```

3.4.2   Empty Sections

Empty sections are allowed at both internal and leaf node levels, enabling the author to express the fact that there is in some particular instance, no data for an attribute, while still showing that the attribute itself is expected to exist in the underlying information model. An empty section looks as follows:

```
address = <>          -- person's address
```

Nested empty sections can be used.

Note: within this document, empty sections are shown in many places to represent fully populated data, which would of course require much more space.

> **Empty sections** can appear anywhere.

3.4.3   Container Objects

The syntax described so far allows an instance of an arbitrarily large object to be expressed, but does not yet allow for attributes of container types such as lists, sets and hash tables, i.e. items whose type in an underlying reference model is something like `attr:List<Type>`, `attr:Set<Type>` or `attr: Hash<ValueType, KeyType>`. There are two ways instance data of such container objects can be expressed in dADL. The first is to use a list style literal value, where the "list nature" of the data is expressed within the manifest value itself, as in the following examples.

```
fruits = <"pear", "cumquat", "peach">
some_primes = <1, 2, 3, 5>
```

See Lists of Built-in Types on page 34 for the complete description of list leaf types. This approach is fine for leaf data. However for containers holding non-primitive values, including more container objects, a different syntax is needed. Consider by way of example that an instance of the container `List<Person>` could be expressed as follows.

```
-- WARNING: THIS IS NOT VALID dADL
people = <
    <name = <> date_of_birth = <> sex = <> interests = <>>
    <name = <> date_of_birth = <> sex = <> interests = <>>
    -- etc
>
```

Here, "anonymous" blocks of data are repeated inside the outer block. However, this makes the data hard to read, and does not provide an easy way of constructing paths to the contained items. A better syntax becomes more obvious when we consider that members of container objects in their computable form are nearly always accessed by a method such as `member(i)`, `item[i]` or just plain `[i]`, in the case of array access in the C-based languages. dADL opts for the array-style syntax, known in dADL as container member *keys*. No attribute name is explicitly given (see Syntax Alternatives on page 44 for further discussion of this choice); any primitive comparable value is allowed as the key, rather than just integers used in C-style array access. Further, if integers are used, it is not assumed that they dictate ordinal indexing, i.e. it is possible to use a series of keys `[2]`, `[4]`, `[8]` etc. The following example shows one version of the above container in valid dADL:

```
people = <
    [1] = <name = <> birth_date = <> interests = <>>
    [2] = <name = <> birth_date = <> interests = <>>
    [3] = <name = <> birth_date = <> interests = <>>
>
```

Strings and dates may also be used. Keys are coloured blue in the this specification in order to distinguish the run-time status of key values from the design-time status of class and attribute names. The following example shows the use of string values as keys for the contained items.

```
people = <
    ["akmal:1975-04-22"] = <name = <> birth_date = <> interests = <>>
    ["akmal:1962-02-11"] = <name = <> birth_date = <> interests = <>>
    ["gianni:1978-11-30"] = <name = <> birth_date = <> interests = <>>
>
```

The syntax for primitive values used as keys follows exactly the same syntax described below for data of primitive types. It is convenient in some cases to construct key values from one or more of the

values of the contained items, in the same way as relational database keys are constructed from suffi-cient field values to guarantee uniqueness. However, they need not be - they may be independent of the contained data, as in the case of hash tables, where the keys are part of the hash table structure, or equally, they may simply be integer index values, as in the 'locations' attribute in the 'school_schedule' structure shown below.

Container structures can appear anywhere in an overall instance structure, allowing complex data such as the following to be expressed in a readable way.

```
school_schedule = <
    lesson_times = <08:30:00, 09:30:00, 10:30:00, ...>

    locations = <
        [1] = <"under the big plane tree">
        [2] = <"under the north arch">
        [3] = <"in a garden">
    >

    subjects = <
        ["philosophy:plato"] = < -- note construction of key
            name = <"philosophy">
            teacher = <"plato">
            topics = <"meta-physics", "natural science">
            weighting = <76%>
        >
        ["philosophy:kant"] = <
            name = <"philosophy">
            teacher = <"kant">
            topics = <"meaning and reason", "meta-physics", "ethics">
            weighting = <80%>
        >
        ["art"] = <
            name = <"art">
            teacher = <"goya">
            topics = <"technique", "portraiture", "satire">
            weighting = <78%>
        >
    >
```

Container instances are expressed using repetitions of a block introduced by a *key*, in the form of a primitive value in brackets i.e. '[]'.

The example above conforms directly to the object-oriented type specification (given in a pascal-like syntax):

```
class SCHEDULE
    lesson_times: List<Time>
    locations: List<String>
    subjects: List<SUBJECT> -- or it could be Hash<SUBJECT>
end

class SUBJECT
    name: String
    teacher: String
    topics: List<String>
    weighting: Real
end
```

Other class specifications corresponding to the same data are possible, but will all be isomorphic to the above.

How key values relate to a particular object structure depends on the class model of objects being created due to a dADL parsing process. It is possible to write a parser which makes reasonable inferences from a class model whose instances are represented as dADL text; it is also possible to include explicit typing information in the dADL itself (see Adding Type Information below).

3.4.3.1   Paths

Paths through container objects are formed in the same way as paths in other structured data, with the addition of the key, to ensure uniqueness. The key is included syntactically enclosed in brackets, in a similar way to how keys are included in Xpath expressions. Paths through containers in the above example include the following:

```
/school_schedule/locations[1]/          -- path to "under the big..."
/school_schedule/subjects["philosophy:kant"]/      -- path to "kant"
```

3.4.4   Nested Container Objects

In some cases the data of interest are instances of nested container types, such as `List<List<Message>>` (a list of Message lists) or `Hash<List<Integer>, String>` (a hash of integer lists keyed by strings). The dADL syntax for such structures follows directly from the syntax for a single container object. The following example shows an instance of the type `List<List<String>>` expressed in dADL syntax.

```
list_of_string_lists = <
    [1] = <
        [1] = <"first string in first list">
        [2] = <"second string in first list">
    >
    [2] = <
        [1] = <"first string in second list">
        [2] = <"second string in second list">
        [3] = <"third string in second list">
    >
    [3] = <
        [1] = <"only string in third list">
    >
>
```

3.4.4.1   Paths

The paths of the above example are as follows:

```
/list_of_string_lists[1]/[1]
/list_of_string_lists[1]/[2]
/list_of_string_lists[2]/[1]
etc
```

3.4.5   Adding Type Information

In many cases, dADL data is of a simple structure, very regular, and highly repetitive, such as the expression of simple demographic data. In such cases, it is preferable to express as little as possible about the implied reference model of the data (i.e. the object or relational model to which it conforms), since various software components want to use the data, and use it in different ways. However, there are also cases where the data is highly complex, and more model information is needed to help software parse it,. Examples include large design databases such as for aircraft, and health records. Typing information is added to instance data using a syntactical addition inspired by the

(type) casting operator of the C language, whose meaning is approximately: force the type of the result of the following expression to be type. In dADL typing is therefore done by including the type name in parentheses after the '=' sign, as in the following example.

```
destinations = <
    ["seville"] = (TOURIST_DESTINATION) <
        profile = (DESTINATION_PROFILE) <>
        hotels = <
            ["gran sevilla"] = (HISTORIC_HOTEL) <>
            ["sofitel"] = (LUXURY_HOTEL) <>
            ["hotel real"] = (PENSION) <>
        >
        attractions = <
            ["la corrida"] = (ATTRACTION) <>
            ["Alcázar"] = (HISTORIC_SITE) <>
        >
    >
>
```

Note that in the above, no type identifiers are included after the "hotels" and "attractions" attributes, and it is up to the processing software to infer the correct types (usually easy - it will be pre-determined by an information model). However, the complete typing information can be included, as follows.

```
hotels = (List<HOTEL>) <
    ["gran sevilla"] = (HISTORIC_HOTEL) <>
>
```

This illustrates the use of generic, or "template" type identifiers, expressed in the standard UML syntax, using angle brackets. Any number of template arguments and any level of nesting is allowed, as in the UML. There is a small risk of visual confusion between the template type delimiters and the standard dADL block delimiters, but technically there can never be any confusion, because only type names (first letter capitalised) may appear inside template delimiters, while only attribute names (first letter lower case) can appear after a dADL block delimiter.

Type identifiers can also include namespace information, which is necessary when same-named types appear in different packages of a model. Namespaces are included by prepending package names, separated by the '.' character, in the same way as in most programming languages, as in the qualified type names **org.openehr.rm.ehr.content.ENTRY** and **Core.Abstractions.Relationships.Relationship**.

> **Type Information** can be included optionally on any node immediately before the opening '<' of any block, in the form of a UML-style type identifier in parentheses. Dot-separated namespace identifiers and template parameters may be used.

3.4.6   Associations and Shared Objects

All of the facilities described so far allow any object-oriented data to be faithfully expressed in a formal, systematic way which is both machine- and human-readable, and allow any node in the data to be addressed using an Xpath-style path. The availability of reliable paths allows not only the representation of single 'business objects', which are the equivalent of UML *aggregation* (and *composition*) hierarchies, but also the representation of *associations* between objects, and by extension, shared objects.

Consider that in the example above, 'hotel' objects may be shared objects, referred to by assocation. This can be expressed as follows.

```
destinations = <
```

```
        ["seville"] = <
            hotels = <
                ["gran sevilla"] = </hotels["gran sevilla"]>
                ["sofitel"] = </hotels["sofitel"]>
                ["hotel real"] = </hotels["hotel real"]>
            >
        >
    >

    bookings = <
        ["seville:0134"] = <
            customer_id = <"0134">
            period = <...>
            hotel = </hotels["sofitel"]>
        >
    >

    hotels = <
        ["gran sevilla"] = (HISTORIC_HOTEL) <>
        ["sofitel"] = (LUXURY_HOTEL) <>
        ["hotel real"] = (PENSION) <>
    >
```

Associations are expressed via the use of fully qualified paths as the data for a attribute. In this example, there are references from a list of destinations, and from a booking list, to the same hotel object. If type information is included, it should go in the declarations of the relevant objects; type declarations can also be used before path references, which might be useful if the association type is an ancestor type (i.e. more general type) of the type of the actual object being referred to.

Data in other dADL documents can be referred to using the URI syntax to locate the document, with the internal path included as described above.

> **Shared objects** are referenced using paths. Objects in other dADL documents can be referred to using normal URIs whose path section conforms to dADL path syntax.

3.4.6.1   Paths

The path set from the above example is as follows:

```
/destinations["seville"]/hotels["gran sevilla"]/
/destinations["seville"]/hotels["sofitel"]/
/destinations["seville"]/hotels["hotel real"]/

/bookings["seville:0134"]/customer_id/
/bookings["seville:0134"]/period/
/bookings["seville:0134"]/hotel/

/hotels["sofitel"]/
/hotels["hotel real"]/
/hotels["gran sevilla"]/
```

3.5   Leaf Data - Built-in Types

All dADL data eventually devolve to instances of the primitive types `String`, `Integer`, `Real`, `Double`, `String`, `Character`, various date/time types, lists or intervals of these types, and a few special types. dADL does not use type or attribute names for instances of primitive types, only manifest values, making it possible to assume as little as possible about type names and structures of the primitive

types. In all the following examples, the manifest data values are assumed to appear immediately inside a leaf pair of angle brackets, i.e.

```
some_attribute = <manifest value here>
```

3.5.1   Primitive Types

3.5.1.1   Character Data

Characters are shown in a number of ways. In the literal form, a character is shown in single quotes, as follows:

```
'a'
```

Special characters are expressed using the ISO 10646 or XML special character codes as described above. Examples:

```
'&ohgr;'      -- greek omega
```

All characters are case-sensitive, i.e. 'a' is distinct from 'A'.

3.5.1.2   String Data

All strings are enclosed in double quotes, as follows:

```
"this is a string"
```

Quotes are encoded using ISO/IEC 10646 codes, e.g. :

```
"this is a much longer string, what one might call a &quot;phrase&quot;."
```

Line extension of strings is done simply by including returns in the string. The exact contents of the string are computed as being the characters between the double quote characters, with the removal of white space leaders up to the left-most character of the first line of the string. This has the effect of allowing the inclusion of multi-line strings in dADL texts, in their most natural human-readable form, e.g.:

```
text = <"And now the STORM-BLAST came, and he
        Was tyrannous and strong :
        He struck with his o'ertaking wings,
        And chased us south along.">
```

String data can be used to contain almost any other kind of data, which is intended to be parsed as some other formalism. Special characters are expressed using the ISO 10646 or XML special character codes within single quotes. ISO codes are mnemonic, and follow the pattern `&aaaa;`, while XML codes are hexadecimal and follow the pattern `&#xHHHH;`, where `H` stands for a hexadecimal digit. An example is:

```
"a &isin; A"      -- prints as: a ∈ A
```

All strings are case-sensitive, i.e. 'word' is distinct from 'Word'.

3.5.1.3   Integer Data

Integers are represented simply as numbers, e.g.:

```
25
300000
29e6
```

Commas or periods for breaking long numbers are not allowed, since they confuse the use of commas used to denote list items (see section 3.5.4 below).

3.5.1.4   Real Data

Real numbers are assumed whenever a decimal is detected in a number, e.g.:

```
25.0
```

```
3.1415926
6.023e23
```

Commas or periods for breaking long numbers are not allowed. Only periods may be used to separate the decimal part of a number; unfortunately, the European use of the comma for this purpose conflicts with the use of the comma to distinguish list items (see section 3.5.4 below).

3.5.1.5   Boolean Data

Boolean values can be indicated by the following values (case-insensitive):

```
True
False
```

3.5.1.6   Dates and Times

Complete Date/Times

In dADL, full and partial dates, times and durations can be expressed. All full dates, times and durations are expressed using a subset of ISO8601. The Support IM provides a full explanation of the ISO8601 semantics supported in openEHR.

In dADL, the use of ISO 8601 allows extended form only (i.e. ':' and '-' must be used). The ISO 8601 method of representing partial dates consisting of a single year number, and partial times consisting of hours only are not supported, since they are ambiguous. See below for partial forms.

Patterns for complete dates and times in dADL include the following:

```
yyyy-MM-dd                          -- a date
hh:mm:ss[,sss][Z|+/-hhmm]           -- a time with optional seconds
yyyy-MM-ddThh:mm:ss[,sss][Z]        -- a date/time
```

where:

```
yyyy    = four-digit year
MM      = month in year
dd      = day in month
hh      = hour in 24 hour clock
mm      = minutes
ss,sss  = seconds, incuding fractional part
Z       = the timezone in the form of a '+' or '-' followed by 4 digits
          indicating the hour offset, e.g. +0930, or else the literal 'Z'
          indicating +0000 (the Greenwich meridian).
```

Durations are expressed using a string which starts with 'P', and is followed by a list of periods, each appended by a single letter designator: 'Y' for years, "M" for months, 'W' for weeks, 'D' for days, 'H' for hours, 'M' for minutes, and 'S' for seconds. The literal 'T' separates the YMWD part from the HMS part, ensuring that months and minutes can be distinguished. Examples of date/time data include:

```
1919-01-23                 -- birthdate of Django Reinhardt
16:35:04,5                 -- rise of Venus in Sydney on 24 Jul 2003
2001-05-12T07:35:20+1000   -- timestamp on an email received from Australia
P22D4TH15M0S               -- period of 22 days, 4 hours, 15 minutes
```

Partial Date/Times

Two ways of expressing partial (i.e. incomplete) date/times are supported in dADL. The ISO 8601 incomplete formats are supported in extended form only (i.e. with '-' and ':' separators) for all patterns that are unambiguous on their own. Dates consisting of only the year, and times consisting of only the hour are not supported, since both of these syntactically look like integers. The supported ISO 8601 patterns are as follows:

```
yyyy-MM                                 -- a date with no days
hh:mm                                   -- a time with no seconds
yyyy-MM-ddThh:mm                        -- a date/time with no seconds
yyyy-MM-ddThh                           -- a date/time, no minutes or seconds
```

To deal with the limitations of ISO 8601 partial patterns in a context-free parsing environment, a second form of pattern is supported in dADL, based on ISO 8601. In this form, '?' characters are substituted for missing digits. Valid partial dates follow the patterns:

```
yyyy-MM-??                  -- date with unknown day in month
yyyy-??-??                  -- date with unknown month and day
```

Valid partial times follow the patterns:

```
hh:mm:??                    -- time with unknown seconds
hh:??:??                    -- time with unknown minutes and seconds
```

Valid date/times follow the patterns:

```
yyyy-MM-ddThh:mm:??         -- date/time with unknown seconds
yyyy-MM-ddThh:??:??         -- date/time with unknown minutes and seconds
yyyy-MM-ddT??:??:??         -- date/time with unknown time
yyyy-MM-??T??:??:??         -- date/time with unknown day and time
yyyy-??-??T??:??:??         -- date/time with unknown month, day and time
```

3.5.2   Intervals of Ordered Primitive Types

Intervals of any ordered primitive type, i.e., Integer, Real, Date, Time, Date_time and Duration, can be stated using the following uniform syntax, where N, M are instances of any of the ordered types:

`	N..M	`	the inclusive range where N and M are integers, or the infinity indicator;
`	<N	`	less than N;
`	>N	`	greater than N;
`	>=N	`	greater than or equal to N;
`	<=N	`	less than or equal to N;
`	N +/-M	`	interval of N ± M.
`!= N`	does not equal N;

The allowable values for N and M include any value in the range of the relevant type, as well as:

```
infinity

-infinity
```

`*`   equivalent to `infinity`

Examples of this syntax include:

```
|0..5|                      -- integer interval
|0.0..1000.0|               -- real interval
|08:02..09:10|              -- interval of time
|>= 1939-02-01|             -- open-ended interval of dates
|5.0 +/-0.5|                -- 4.5 - 5.5
|>=0|                       -- >= 0
|0..infinity|               -- 0 - infinity (i.e. >= 0)
```

3.5.3   Other Built-in Types

3.5.3.1   URIs

URI can be expressed as dADL data in the usual way found on the web, and follow the standard syntax from http://www.ietf.org/rfc/rfc2396.txt. In the context of dADL, no quotes or inverted

commas are needed; and neither spaces nor angle brackets are allowed; both have to be quoted e.g. %20 means a space (as per the URI standard). Examples of URIs in dADL:

```
http://archetypes.are.us/home.html
ftp://get.this.file.com#section_5
http://www.mozilla.org/products/firefox/upgrade/?application=thunderbird
```

Encoding of special characters in URIs follows the IETF RFC 3986, as described under Reserved Character Encoding on page 23.

3.5.3.2   Coded Terms

Coded terms are ubiquitous in medical and clinical information, and are likely to become so in most other industries, as ontologically-based information systems and the 'semantic web' emerge. The logical structure of a coded term is simple: it consists of an identifier of a terminology, and an identifier of a code within that terminology. The dADL string representation is as follows:

```
[terminology_id::code]
```

Typical examples from clinical data:

```
[icd10AM::F60.1]           -- from ICD10AM
[snomed-ct::2004950]       -- from snomed-ct
[snomed-ct(3.1)::2004950]  -- from snomed-ct v 3.1
```

3.5.4   Lists of Built-in Types

Data of any primitive type can occur singly or in lists, which are shown as comma-separated lists of item, all of the same type, such as in the following examples:

```
"cyan", "magenta", "yellow", "black"  -- printer's colours
1, 1, 2, 3, 5                          -- first 5 fibonacci numbers
08:02, 08:35, 09:10                    -- set of train times
```

No assumption is made in the syntax about whether a list represents a set, a list or some other kind of sequence - such semantics must be taken from an underlying information model.

Lists which happen to have only one datum are indicated by using a comma followed by a list continuation marker of three dots, i.e. "...", e.g.:

```
"en", ...        -- languages
"icd10", ...     -- terminologies
[at0200], ...
```

White space may be freely used or avoided in lists, i.e. the following two lists are identical:

```
1,1,2,3
1, 1, 2,3
```

3.6   Plug-in Syntaxes

Using the dADL syntax, any object structure can be serialised. In some cases, the requirement is to express some part of the structure in an abstract syntax, rather than in the more literal seriliased object form of dADL. dADL provides for this possibility by allowing the value of any object (i.e. what appears between any matching pair of <> delimiters) to be expressed in some other syntax, known as a "plug-in" syntax. Plug-in syntaxes are indicated in dADL in a similar way as typed objects, i.e. by the use of the syntax type in parentheses preceding the <> block. For a plug-in section, the <> delimiters are modified to <# #>, to allow for easier parser design, and easier recognition of such blocks by human readers. The general form is as follows:

```
attr_name = (syntax) <#
    ...
```

```
    #>
```

The following example illustrates a cADL plug-in section in an archetype, which it itself a dADL document:

```
definition = (cadl) <#
    ENTRY[at0000] ∈ {              -- blood pressure measurement
        name ∈ {                   -- any synonym of BP
            CODED_TEXT ∈ {
                code ∈ {
                    CODE_PHRASE ∈ {[ac0001]}
                }
            }
        }
    }
#>
```

Clearly, many plug-in syntaxes might one day be used within dADL data; there is no guarantee that every dADL parser will support them. The general approach to parsing should be to use plug-in parsers, i.e. to obtain a parser for a plug-in syntax that can be built into the existing parser framework.

3.7   Expression of dADL in XML

The dADL syntax maps quite easily to XML instance. It is important to realise that people using XML often develop different mappings for object-oriented data, due to the fact that XML does not have systematic object-oriented semantics. This is particularly the case where containers such as lists and sets such as 'employees: List<Person>' are mapped to XML; many implementors have to invent additional tags such as 'employee' to make the mapping appear visually correct. The particular mapping chosen here is designed to be a faithful reflection of the semantics of the object-oriented data, and does not try take into account visual aesthetics of the XML. The result is that Xpath expressions are the same for dADL and XML, and also correspond to what one would expect based on an underlying object model.

The main elements of the mapping are as follows.

Single Attributes

Single attribute nodes map to tagged nodes of the same name.

Container Attributes

Container attribute nodes map to a series of tagged nodes of the same name, each with the XML attribute 'id' set to the dADL key. For example, the dADL:

```
subjects = <
    ["philosophy:plato"] = <
        name = <"philosophy">
    >
    ["philosophy:kant"] = <
        name = <"philosophy">
    >
>
```

maps to the XML:

```
<subjects id="philosophy:plato">
        <name>
            philosophy
        </name>
</subjects>
```

```
<subjects id="philosophy:kant">
    <name>
        philosophy
    </name>
</subjects>
```

This guarantees that the path `subjects[@id="philosophy:plato"]/name` navigates to the same element in both dADL and the XML.

Nested Container Attributes

Nested container attribute nodes map to a series of tagged nodes of the same name, each with the XML attribute 'id' set to the dADL key. For example, consider an object structure defined by the signature `countries:Hash<Hash<Hotel,String>,String>`. An instance of this in dADL looks as follows:

```
countries = <
    ["spain"] = <
        ["hotels"] = <...>
        ["attractions"] = <...>
    >
    ["egypt"] = <
        ["hotels"] = <...>
        ["attractions"] = <...>
    >
>
```

can be mapped to the XML in which the synthesised element tag "_items" and the attribute "key" are used:

```
<countries key="spain">
    <_items key="hotels">
        ...
    </_items>
    <_items key="attractions">
        ...
    </_items>
<countries>
</countries key="eqypt">
    <_items id="hotels">
        ...
    </_items>
    <_items key="attractions">
        ...
    </_items>
</countries>
```

In this case, the dADL path `countries["spain"]/["hotels"]` will be transformed to the Xpath `countries[@key="spain"]/_items[@key="hotels"]` in order to navigate to the same element.

Type Names

Type names map to XML 'type' attributes e.g. the dADL:

```
destinations = <
    ["seville"] = (TOURIST_DESTINATION) <
        profile = (DESTINATION_PROFILE) <>
        hotels = <
            ["gran sevilla"] = (HISTORIC_HOTEL) <>
        >
    >
```

```
>
```

maps to:

```
<destinations id="seville" xsi:type=TOURIST_DESTINATION>
        <profile xsi:type=DESTINATION_PROFILE>
          ...
        </profile>
        <hotels id="gran sevilla" xsi:type=HISTORIC_HOTEL>
          ...
        </hotels>
    >
>
```

3.8   Syntax Specification

The dADL grammar is available as an HTML_document. This grammar is implemented and tested using lex (.l file) and yacc (.y file) specifications for in the Eiffel programming environment. The current release of these files is available at http://svn.openehr.org/ref_impl_eiffel/TRUNK/libraries/common_libs/src/structures/syntax/dadl/parser/. The .l and .y files can easily be converted for use in another yacc/lex-based programming environment.

3.8.1   Grammar

The following provides the dADL parser production rules (yacc specification) as of revision 36 of the Eiffel reference implementation repository (http://svn.openehr.org/ref_impl_eiffel).

```
input:
  attr_vals
| complex_object_block
| error

attr_vals:
  attr_val
| attr_vals attr_val
| attr_vals ; attr_val

attr_val:
  attr_id SYM_EQ object_block

attr_id:
  V_ATTRIBUTE_IDENTIFIER
| V_ATTRIBUTE_IDENTIFIER error

object_block:
  complex_object_block
| primitive_object_block
| plugin_object_block

plugin_object_block:
  V_PLUGIN_SYNTAX_TYPE V_PLUGIN_BLOCK

complex_object_block:
  single_attr_object_block
| multiple_attr_object_block

multiple_attr_object_block:
```

```
      untyped_multiple_attr_object_block
    | V_TYPE_IDENTIFIER untyped_multiple_attr_object_block
```

untyped_multiple_attr_object_block:
```
    multiple_attr_object_block_head keyed_objects SYM_END_DBLOCK
```

multiple_attr_object_block_head:
```
    SYM_START_DBLOCK
```

keyed_objects:
```
    keyed_object
    | keyed_objects keyed_object
```

keyed_object:
```
    object_key SYM_EQ object_block
```

object_key:
```
    [ simple_value ]
```

single_attr_object_block:
```
    untyped_single_attr_object_block
    | V_TYPE_IDENTIFIER untyped_single_attr_object_block
```

untyped_single_attr_object_block:
```
    single_attr_object_complex_head SYM_END_DBLOCK
    | single_attr_object_complex_head attr_vals SYM_END_DBLOCK
```

single_attr_object_complex_head:
```
    SYM_START_DBLOCK
```

primitive_object_block:
```
    untyped_primitive_object_block
    | V_TYPE_IDENTIFIER untyped_primitive_object_block
```

untyped_primitive_object_block:
```
    SYM_START_DBLOCK primitive_object_value SYM_END_DBLOCK
```

primitive_object_value:
```
    simple_value
    | simple_list_value
    | simple_interval_value
    | term_code
    | term_code_list_value
```

simple_value:
```
    string_value
    | integer_value
    | real_value
    | boolean_value
    | character_value
    | date_value
    | time_value
    | date_time_value
    | duration_value
    | uri_value
```

simple_list_value:

```
  string_list_value
| integer_list_value
| real_list_value
| boolean_list_value
| character_list_value
| date_list_value
| time_list_value
| date_time_list_value
| duration_list_value
```

simple_interval_value:
```
  integer_interval_value
| real_interval_value
| date_interval_value
| time_interval_value
| date_time_interval_value
| duration_interval_value
```

string_value:
```
  V_STRING
```

string_list_value:
```
  V_STRING , V_STRING
| string_list_value , V_STRING
| V_STRING , SYM_LIST_CONTINUE
```

integer_value:
```
  V_INTEGER
| + V_INTEGER
| - V_INTEGER
```

integer_list_value:
```
  integer_value , integer_value
| integer_list_value , integer_value
| integer_value , SYM_LIST_CONTINUE
```

integer_interval_value:
```
  SYM_INTERVAL_DELIM integer_value SYM_ELLIPSIS integer_value
SYM_INTERVAL_DELIM
| SYM_INTERVAL_DELIM SYM_LT integer_value SYM_INTERVAL_DELIM
| SYM_INTERVAL_DELIM SYM_LE integer_value SYM_INTERVAL_DELIM
| SYM_INTERVAL_DELIM SYM_GT integer_value SYM_INTERVAL_DELIM
| SYM_INTERVAL_DELIM SYM_GE integer_value SYM_INTERVAL_DELIM
| SYM_INTERVAL_DELIM integer_value SYM_INTERVAL_DELIM
```

real_value:
```
  V_REAL
| + V_REAL
| - V_REAL
```

real_list_value:
```
  real_value , real_value
| real_list_value , real_value
| real_value , SYM_LIST_CONTINUE
```

real_interval_value:
```
  SYM_INTERVAL_DELIM real_value SYM_ELLIPSIS real_value SYM_INTERVAL_DELIM
```

```
  | SYM_INTERVAL_DELIM SYM_LT real_value SYM_INTERVAL_DELIM
  | SYM_INTERVAL_DELIM SYM_LE real_value SYM_INTERVAL_DELIM
  | SYM_INTERVAL_DELIM SYM_GT real_value SYM_INTERVAL_DELIM
  | SYM_INTERVAL_DELIM SYM_GE real_value SYM_INTERVAL_DELIM
  | SYM_INTERVAL_DELIM real_value SYM_INTERVAL_DELIM

  boolean_value:
    SYM_TRUE
  | SYM_FALSE

  boolean_list_value:
    boolean_value , boolean_value
  | boolean_list_value , boolean_value
  | boolean_value , SYM_LIST_CONTINUE

  character_value:
    V_CHARACTER

  character_list_value:
    character_value , character_value
  | character_list_value , character_value
  | character_value , SYM_LIST_CONTINUE

  date_value:
    V_ISO8601_EXTENDED_DATE

  date_list_value:
    date_value , date_value
  | date_list_value , date_value
  | date_value , SYM_LIST_CONTINUE

  date_interval_value:
    SYM_INTERVAL_DELIM date_value SYM_ELLIPSIS date_value SYM_INTERVAL_DELIM
  | SYM_INTERVAL_DELIM SYM_LT date_value SYM_INTERVAL_DELIM
  | SYM_INTERVAL_DELIM SYM_LE date_value SYM_INTERVAL_DELIM
  | SYM_INTERVAL_DELIM SYM_GT date_value SYM_INTERVAL_DELIM
  | SYM_INTERVAL_DELIM SYM_GE date_value SYM_INTERVAL_DELIM
  | SYM_INTERVAL_DELIM date_value SYM_INTERVAL_DELIM

  time_value:
    V_ISO8601_EXTENDED_TIME

  time_list_value:
    time_value , time_value
  | time_list_value , time_value
  | time_value , SYM_LIST_CONTINUE

  time_interval_value:
    SYM_INTERVAL_DELIM time_value SYM_ELLIPSIS time_value SYM_INTERVAL_DELIM
  | SYM_INTERVAL_DELIM SYM_LT time_value SYM_INTERVAL_DELIM
  | SYM_INTERVAL_DELIM SYM_LE time_value SYM_INTERVAL_DELIM
  | SYM_INTERVAL_DELIM SYM_GT time_value SYM_INTERVAL_DELIM
  | SYM_INTERVAL_DELIM SYM_GE time_value SYM_INTERVAL_DELIM
  | SYM_INTERVAL_DELIM time_value SYM_INTERVAL_DELIM

  date_time_value:
    V_ISO8601_EXTENDED_DATE_TIME
```

```
date_time_list_value:
   date_time_value , date_time_value
| date_time_list_value , date_time_value
| date_time_value , SYM_LIST_CONTINUE


date_time_interval_value:
   SYM_INTERVAL_DELIM date_time_value SYM_ELLIPSIS date_time_value
SYM_INTERVAL_DELIM
| SYM_INTERVAL_DELIM SYM_LT date_time_value SYM_INTERVAL_DELIM
| SYM_INTERVAL_DELIM SYM_LE date_time_value SYM_INTERVAL_DELIM
| SYM_INTERVAL_DELIM SYM_GT date_time_value SYM_INTERVAL_DELIM
| SYM_INTERVAL_DELIM SYM_GE date_time_value SYM_INTERVAL_DELIM
| SYM_INTERVAL_DELIM date_time_value SYM_INTERVAL_DELIM


duration_value:
   V_ISO8601_DURATION
| - V_ISO8601_DURATION


duration_list_value:
   duration_value , duration_value
| duration_list_value , duration_value
| duration_value , SYM_LIST_CONTINUE


duration_interval_value:
   SYM_INTERVAL_DELIM duration_value SYM_ELLIPSIS duration_value
SYM_INTERVAL_DELIM
| SYM_INTERVAL_DELIM SYM_LT duration_value SYM_INTERVAL_DELIM
| SYM_INTERVAL_DELIM SYM_LE duration_value SYM_INTERVAL_DELIM
| SYM_INTERVAL_DELIM SYM_GT duration_value SYM_INTERVAL_DELIM
| SYM_INTERVAL_DELIM SYM_GE duration_value SYM_INTERVAL_DELIM
| SYM_INTERVAL_DELIM duration_value SYM_INTERVAL_DELIM


term_code:
   V_QUALIFIED_TERM_CODE_REF


term_code_list_value:
   term_code , term_code
| term_code_list_value , term_code
| term_code , SYM_LIST_CONTINUE


uri_value:
   V_URI
```

3.8.2   Symbols

The following provides the dADL lexical analyser production rules (lex specification) as of revision 36 of the Eiffel reference implementation repository (http://svn.openehr.org/ref_impl_eiffel):

```
----------/** Separators **/------------------------------------------


[ \t\r]+            -- Ignore separators
\n+                 -- (increment line count)


----------/** comments **/---------------------------------------------


"--".*              -- Ignore comments
```

```
    "--".*\n[ \t\r]*    -- (increment line count)


    ----------/* symbols */ ------------------------------------------
    "-"          -- -> Minus_code
    "+"          -- -> Plus_code
    "*"          -- -> Star_code
    "/"          -- -> Slash_code
    "^"          -- -> Caret_code
    "."          -- -> Dot_code
    ";"          -- -> Semicolon_code
    ","          -- -> Comma_code
    ":"          -- -> Colon_code
    "!"          -- -> Exclamation_code
    "("          -- -> Left_parenthesis_code
    ")"          -- -> Right_parenthesis_code
    "$"          -- -> Dollar_code
    "??"         -- -> SYM_DT_UNKNOWN
    "?"          -- -> Question_mark_code

    "|"          -- -> SYM_INTERVAL_DELIM

    "["          -- -> Left_bracket_code
    "]"          -- -> Right_bracket_code


    "="          -- -> SYM_EQ

    ">="         -- -> SYM_GE
    "<="         -- -> SYM_LE


    "<"          -- -> SYM_LT or SYM_START_DBLOCK
    ">"          -- -> SYM_GT or SYM_END_DBLOCK


    ".."         -- -> SYM_ELLIPSIS
    "..."        -- -> SYM_LIST_CONTINUE

    ----------/* keywords */ ------------------------------------------

    [Tt][Rr][Uu][Ee]         -- -> SYM_TRUE

    [Ff][Aa][Ll][Ss][Ee]     -- -> SYM_FALSE

    [Ii][Nn][Ff][Ii][Nn][Ii][Tt][Yy]      -- -> SYM_INFINITY

    ----------/* V_URI */ ---------------------------------------------
    [a-z]+:\/\/[^> ]*

    ----------/* V_QUALIFIED_TERM_CODE_REF form [ICD10AM(1998)::F23] */ -----
    \[[a-zA-Z0-9()._\-]+::[a-zA-Z0-9._\-]+\]

    ----------/* V_LOCAL_TERM_CODE_REF */ -----------------------------
    \[[a-zA-Z0-9][a-zA-Z0-9._\-]*\]

    ----------/* V_LOCAL_CODE */ --------------------------------------
    a[ct][0-9.]+

    ----------/* V_ISO8601_DATE_TIME YYYY-MM-DDThh:mm:ss[,sss][Z|+/-nnnn] */
    [0-9]{4}-[0-1][0-9]-[0-3][0-9]T[0-2][0-9]:[0-6][0-9]:[0-6][0-9](,[0-9]+)?(Z|[+-][0-9]{4})?  |
```

```
[0-9]{4}-[0-1][0-9]-[0-3][0-9]T[0-2][0-9]:[0-6][0-9](Z|[+-][0-9]{4})? |
[0-9]{4}-[0-1][0-9]-[0-3][0-9]T[0-2][0-9](Z|[+-][0-9]{4})?

----------/* V_ISO8601_TIME Thh:mm:ss[,sss][Z|+/-nnnn] */ ------------
T[0-2][0-9]:[0-6][0-9]:[0-6][0-9](,[0-9]+)?(Z|[+-][0-9]{4})? |
T[0-2][0-9]:[0-6][0-9](Z|[+-][0-9]{4})? |
T[0-2][0-9](Z|[+-][0-9]{4})?

----------/* V_ISO8601_DATE YYYY-MM-DD */ ---------------------------------
[0-9]{4}-[0-1][0-9]-[0-3][0-9] |
[0-9]{4}-[0-1][0-9] {
-- [0-9]{4}  -- can't handle year only, since it will be confused with plain integer

----------/* V_ISO8601_DURATION */ ------------------------------------
P[0-9]+[wW]|
P([0-9]+[yY])?([0-9]+[mM])?([0-9]+[dD])?T([0-9]+[hH])?([0-9]+[mM])?([0-9]+[sS])? |
P([0-9]+[yY])?([0-9]+[mM])?([0-9]+[dD])?

----------/* V_TYPE_IDENTIFIER */ -------------------------------------
[A-Z][a-zA-Z0-9_]*

----------/* V_ATTRIBUTE_IDENTIFIER */ -------------------------------------
[a-z][a-zA-Z0-9_]*

----------/* CADL Blocks */ -----------------------------------------
\{[^{}]*                -- beginning of CADL block
<IN_CADL_BLOCK>\{[^{}]*  -- got an open brace
<IN_CADL_BLOCK>[^{}]*\}  -- got a close brace

----------/* V_INTEGER */ -----------------------------------------
[0-9]+ |
[0-9]+[eE][+-]?[0-9]+

----------/* V_REAL */ -----------------------------------------
[0-9]+\.[0-9]+|
[0-9]+\.[0-9]+[eE][+-]?[0-9]+

----------/* V_STRING */ -----------------------------------------
\"[^\\\n"]*\"

\"[^\\\n"]*{     -- beginning of a multi-line string
<IN_STR>\\\\     -- match escaped backslash, i.e. \\ -> \
<IN_STR>\\\"     -- match escaped double quote, i.e. \" -> "
<IN_STR>&[a-zA-Z][a-zA-Z0-9_]+;-- match ISO special character pattern
&char_name;
<IN_STR>&#x[a-fA-F0-9]{4}; -- match W3C XML special character pattern
&#xHHHH;
<IN_STR>[^\\\n"]+     -- match any other characters
<IN_STR>\\\n[ \t\r]*  -- match LF in line
<IN_STR>[^\\\n"]*\"   -- match final end of string
<IN_STR>.|\n|
<IN_STR><<EOF>>      -- unclosed String -> ERR_STRING

----------/* V_CHARACTER */ -----------------------------------------
\'[^\\\n']\'
\'\\n\
\'\\r\
\'\\t\
```

```
\'\\f\
\'\\'\
\'\\\\
\'\\\"\
\'\\[0-9]+\'
\'&[a-zA-Z0-9_]+;\'  -- match ISO special character pattern &char_name;
\'&#x[a-fA-F0-9]{4};\'  -- match W3C XML special character pattern &#xHHHH;
\'.{1,2}  |
\'\\[0-9]+(\/)?  -- invalid character -> ERR_CHARACTER
```

3.9   Syntax Alternatives

WARNING:the syntax in this section is not part of dADL

3.9.1   Container Attributes

A reasonable alternative to the syntax described above for nested container objects would have been to use an arbitrary member attribute name, such as "items", or perhaps "_items" (in order to indicate to a parser that the attribute name cannot be assumed to correspond to a real property in an object model), as well as the key for each container member, giving syntax like the following:

```
people = <
    _items[1] = <name = <> birth_date = <> interests = <>>
    _items[2] = <name = <> birth_date = <> interests = <>>
    _items[3] = <name = <> birth_date = <> interests = <>>
>
```

Additionally, with this alternative, it becomes more obvious how to include the values of other properties of container types, such as ordering, maximum size and so on, e.g.:

```
people = <
    _items[1] = <name = <> birth_date = <> interests = <>>
    _items[2] = <name = <> birth_date = <> interests = <>>
    _items[3] = <name = <> birth_date = <> interests = <>>
    _is_ordered = <True>
    _upper = <200>
>
```

Again, since the names of such properties in any given object technology cannot be assumed, the special underscore form of attribute names is used.

However, we are now led to somewhat clumsy paths, where "_items" will occur very frequently, due to the ubiquity of containers in real data:

```
/people/_items[1]/
/people/_items[2]/
/people/_items[3]/
/people/_is_ordered/
/people/_upper/
```

A compromise which satisfies the need for correct representation of all attributes of container types and the need for brevity and comprehensibility of paths would be to make optional the "_items", but retain other container pseudo-attributes (likely to be much more rarely used), thus:

```
people = <
    [1] = <name = <> birth_date = <> interests = <>>
    [2] = <name = <> birth_date = <> interests = <>>
    [3] = <name = <> birth_date = <> interests = <>>
    _is_ordered = <True>
    _upper = <200>
```

```
    >
```
The above form leads to the following paths:

```
/people[1]/
/people[2]/
/people[3]/
/people/_is_ordered/
/people/_upper/
```

The alternative syntax in this subsection is not currently part of dADL, but could be included in the future, if there was a need to support more precise modelling of container types in dADL. If such support were to be added, it is recommended that the names of the pseudo-attributes ("_item", "_is_ordered" etc) be based on names of appropriate container types from a recognised standard such as OMG UML, OCL or IDL.

4 cADL - Constraint ADL

4.1 Overview

cADL is a syntax which enables constraints on data defined by object-oriented information models to be expressed in archetypes or other knowledge definition formalisms. It is most useful for defining the specific allowable constructions of data whose instances conform to very general object models. cADL is used both at "design time", by authors and/or tools, and at runtime, by computational systems which validate data by comparing it to the appropriate sections of cADL in an archetype. The general appearance of cADL is illustrated by the following example:

```
PERSON[at0000] matches {              -- constraint on PERSON instance
    name matches {                    -- constraint on PERSON.name
        TEXT matches {/.+/}           -- any non-empty string
    }
    addresses cardinality matches {0..*} matches { -- constraint on
        ADDRESS matches {             -- PERSON.addresses
            -- etc --
        }
    }
}
```

Some of the textual keywords in this example can be more efficiently rendered using common mathematical logic symbols. In the following example, the `matches`, `exists` and `implies` keywords have been replaced by appropriate symbols:

```
PERSON[at0000] ∈ {                    -- constraint on PERSON instance
    name ∈ {                          -- constraint on PERSON.name
        TEXT ∈ {/..*/}                -- any non-empty string
    }
    addresses cardinality ∈ {0..*} ∈ { -- constraint on
        ADDRESS ∈ {                   -- PERSON.addresses
            -- etc --
        }
    }
}
```

The full set of equivalences appears below. Raw cADL is stored in the text-based form, to remove any difficulties with representation of symbols, to avoid difficulties of authoring cADL text in basic text editors which do not supply such symbols, and to aid reading in English. However, the symbolic form might be more widely used due to the use of tools, and formatting in HTML and other documentary formats, and may be more comfortable for non-English speakers and those with formal mathematical backgrounds. This document uses both conventions. The use of symbols or text is completely a matter of taste, and no meaning whatsoever is lost by completely ignoring one or other format according to one's personal preference.

In the standard cADL documented in this section, literal leaf values (such as the regular expression /..*/ in the above example) are always constraints on a set of 'standard' widely-accepted primitive types, as described in the dADL section. Other more sophisticated constraint syntax types are described in cADL - Constraint ADL on page 47.

4.2   Basics

4.2.1   Keywords

The following keywords are recognised in cADL:

- `matches, ~matches, is_in, ~is_in`
- `occurrences, existence, cardinality`
- `ordered, unordered, unique`
- `infinity`
- `use_node, allow_archetype`[1]
- `include, exclude`

Symbol equivalents for some of the above are given in the following table.

Textual Rendering	Symbolic Rendering	Meaning
matches, is_in	$\in$	Set membership, "p is in P"
not, ~	~	Negation, "not p"

Keywords are shown in blue in this document.

The `matches` or `is_in` operator deserves special mention, since it is a key operator in cADL. This operator can be understood mathematically as set membership. When it occurs between a name and a block delimited by braces, the meaning is: the set of values allowed for the entity referred to by the name (either an object, or parts of an object - attributes) is specified between the braces. What appears between any matching pair of braces can be thought of as a *specification for a set of values*. Since blocks can be nested, this approach to specifying values can be understood in terms of nested sets, or in terms of a value space for objects of a set of defined types. Thus, in the following example, the `matches` operator links the name of an entity to a linear value space (i.e. a list), consisting of all words ending in "ion".

```
aaa matches {/.*ion[^\s\n\t]/}-- the set of english words ending in 'ion'
```

The following example links the name of a type XXX with a complex multi-dimensional value space.

```
XXX matches {
    aaa matches {                       --
        YYY matches {0..3}              --
    }                                   -- the value space of the
    bbb matches {                       -- and instance of XXX
        ZZZ matches {>1992-12-01}       --
    }                                   --
}
```

The meaning of the constraint structure above is: in data matching the constraints, there is an instance of type XXX whose attribute values recursively match the inner constraints named after those attributes, and so on, to the leaf level.

Occasionally, the `matches` operator needs to be used in the negative, usually at a leaf block. Any of the following can be used to constrain the value space of the attribute aaa to any number except 5:

```
aaa ~matches {5}
```

---

1. was 'use_archetype', which is now deprecated

```
aaa ~is_in {5}
aaa ∉ {5}
```

The choice of whether to use `matches` or `is_in` is a matter of taste and background; those with a mathematical background will probably prefer `is_in`, while those with a data processing background may prefer `matches`.

4.2.2   Comments

In a cADL text, comments satisfy the following rule:

> **comments** are indicated by the characters "--". **Multi-line comments** are achieved using the "--" leader on each line where the comment continues. In this document, comments are shown in brown.

4.2.3   Information Model Identifiers

As with dADL, identifiers from the underlying information model are used for all cADL nodes. Identifiers obey the same rules as in dADL: type names commence with an upper case letter, while attribute and function names commence with a lower case letter. In cADL, type names and any property (i.e. attribute or function) name can be used, whereas in dADL, only type names and attribute names appear.

> A **type name** is any identifier with an initial upper case letter, followed by any combination of letters, digits, and underscores. An **attribute name** is any identifier with an initial lower case letter, followed by any combination of letters, digits and underscores.

Type identifiers are shown in this document in all uppercase, e.g. `PERSON`, while attribute identifiers are shown in all lowercase, e.g. `home_address`. In both cases, underscores are used to represent word breaks. This convention is used to improve the readability of this document, and other conventions may be used, such as the common programmer's mixed-case convention exemplified by `Person` and `homeAddress`. The convention chosen for any particular cADL document should be based on that used in the underlying information model. Identifiers are shown in green in this document.

4.2.4   Node Identifiers

In cADL, an entity in brackets e.g. `[xxxx]` is used to identify "object nodes", i.e. nodes expressing constraints on instances of some type. Object nodes always commence with a type name. Any string may appear within the brackets, depending on how it is used. However, in this document, all node identifiers are of the form of an archetype term identifier, i.e. `[atNNNN]`, e.g. `[at0042]`. Node identifiers are shown in magenta in this document.

4.2.5   Natural Language

cADL is completely independent of all natural languages. The only potential exception is where constraints include literal values from some language, and this is easily and routinely avoided by the use of separate language and terminology definitions, as used in ADL archetypes. However, for the purposes of readability, comments in English have been included in this document to aid the reader. In real cADL documents, comments are generated from the archetype ontology in the local language.

4.3   Structure

cADL constraints are written in a block-structured style, similar to block-structured programming languages like C. A typical block resembles the following (the recurring pattern `/.+/` is a regular expression meaning "non-empty string"):

```
PERSON[at0001] ∈ {
```

```
    name ∈ {
        PERSON_NAME[at0002] ∈ {
            forenames cardinality ∈ {1..*} ∈ {/.+/}
            family_name ∈ {/.+/}
            title ∈ {"Dr", "Miss", "Mrs", "Mr", ...}
        }
    }
    addresses cardinality ∈ {1..*} ∈ {
        LOCATION_ADDRESS[at0003] ∈ {
            street_number existence ∈ {0..1} ∈ {/.+/}
            street_name ∈ {/.+/}
            locality ∈ {/.+/}
            post_code ∈ {/.+/}
            state ∈ {/.+/}
            country ∈ {/.+/}
        }
    }
}
```

In the above, any identifier (shown in green) followed by the ∈ operator (equivalent text keyword: `matches` or `is_in`) followed by an open brace, is the start of a "block", which continues until the closing matching brace (normally visually indented to come under the start of the line at the beginning of the block).

The example above expresses a constraint on an instance of the type PERSON; the constraint is expressed by everything inside the PERSON block. The two blocks at the next level define constraints on properties of PERSON, in this case *name* and *addresses*. Each of these constraints is expressed in turn by the next level containing constraints on further types, and so on. The general structure is therefore a recursive nesting of constraints on types, followed by constraints on properties (of that type), followed by types (being the types of the attribute under which it appears) until leaf nodes are reached.

We use the term "object" block or node to refer to any block introduced by a type name (in this document, in all upper case), while an "attribute" block or node is any block introduced by an attribute identifier (in all lower case in this document), as illustrated below.

```
                                          PERSON ∈ {
                                            name ∈ {
 object block   attribute block   object block   TEXT ∈ {/..*/}
                                            }
                                          }
```

FIGURE 4 Object and Attribute Blocks in cADL

4.3.1 Complex Objects

It may by now be clear that the identifiers in the above could correspond to entities in an object-oriented information model. A UML model compatible with the example above is shown in FIGURE 5. Note that there can easily be more than one model compatible with a given fragment of cADL syntax, and in particular, there may be more properties and classes in the reference model than are mentioned in the cADL constraints. In other words, a cADL text includes constraints *only for those parts of a model which are useful or meaningful to constrain*.

Constraints expressed in cADL cannot be stronger than those from the information model. For example, the PERSON.*family_name* attribute is mandatory in the model in FIGURE 5, so it is not valid to

FIGURE 5 UML Model of PERSON

express a constraint allowing the attribute to be optional. In general, a cADL archetype can only further constrain an existing information model. However, it must be remembered that for very generic models consisting of only a few classes and a lot of optionality, this rule is not so much a limitation as a way of adding meaning to information. Thus, for a demographic information model which has only the types `PARTY` and `PERSON`, one can write cADL which defines the concepts of entities such as `COM-PANY`, `EMPLOYEE`, `PROFESSIONAL`, and so on, in terms of constraints on the types available in the information model.

This general approach can be used to express constraints for instances of any information model. An example showing how to express a constraint on the *value* property of an `ELEMENT` class to be a `QUAN-TITY` with a suitable range for expressing blood pressure is as follows:

```
ELEMENT[at0010] matches {    -- diastolic blood pressure
    value matches {
        QUANTITY matches {
            magnitude matches {|0..1000|}
            property matches {"pressure"}
            units matches {"mm[Hg]"}
        }
    }
}
```

4.3.2   Attribute Constraints

In any information model, attributes are either single-valued or multiply-valued, i.e. of a generic container type such as `List<Contact>`.

4.3.2.1   Existence

The only constraint that applies to all attributes is to do with existence. Existence constraints say whether an attribute must exist, and are indicated by "0..1" or "1" markers at line ends in UML diagrams (and often mistakenly referred to as a "cardinality of 1..1"). It is the absence or presence of the cardinality constraint in cADL which indicates that the attribute being constrained is single-valued or a container attribute, respectively. Existence constraints are expressed in cADL as follows:

```
QUANTITY matches {
    units existence matches {0..1} matches {"mm[Hg]"}
}
```

The meaning of an existence constraint is to indicate whether a value - i.e. an object - is mandatory or optional (i.e. obligatory or not) in runtime data for the attribute in question. The above example indicates that a value for the 'units' attribute is optional. The same logic applies whether the attribute is of single or multiple cardinality, i.e. whether it is a container or not. For container attributes, the existence constraint indicates whether the whole container (usually a list or set) is mandatory or not; a further *cardinality* constraint (described below) indicates how many members in the container are allowed.

> An **existence constraint** may be used directly after any attribute identifier, and indicates whether the object to which the attribute refers is mandatory or optional in the data.

Existence is shown using the same constraint language as the rest of the archetype definition. Existence constraints can take the values {0}, {0..0}, {0..1}, {1}, or {1..1}. The first two of these constraints may not seem initially obvious, but may be reasonable in some cases: they say that an attribute must not be present in the particular situation modelled by the archetype. The default existence constraint, if none is shown, is {1..1}.

4.3.3 Single-valued Attributes

Repeated blocks of object constraints of the same class (or its subtypes) can have two possible meanings in cADL, depending on whether the cardinality is present or not in the containing attribute block. With no cardinality, the meaning is that each child object constraint of the attribute in question is a possible alternative for the value of the attribute in the data, as shown in the following example:

```
ELEMENT[at0004] matches {               -- speed limit
    value matches {
       QUANTITY matches {
          magnitude matches {0..55}
          property matches {"velocity"}
          units matches {"mph"}       -- miles per hour
       }
       QUANTITY matches {
          magnitude matches {0..100}
          property matches {"velocity"}
          units matches {"km/h"}      -- km per hour
       }
    }
}
```

Here, the cardinality of the *value* attribute is 1..1 (the default), while the occurrences of both QUANTITY constraints is optional, leading to the result that only one QUANTITY instance can appear in runtime data, and it can match either of the constraints.

> Two or more object blocks introduced by type names appearing after an attribute which is not a container (i.e. for which there is no cardinality constraint) are taken to be **alternative constraints**, only one of which needs to be matched by the data.

Note that there is a more efficient way to express the above example, using domain type extensions. An example is provided in section 10.1.2 on page 99.

4.3.4 Container Attributes

4.3.4.1 Cardinality

Container attributes are indicated in cADL with the *cardinality* constraint. Cardinalities indicate limits on the number of members of instances of container types such as lists and sets. Consider the following example:

```
HISTORY occurrences ∈ {1} ∈ {
    periodic ∈ {False}
    events cardinality ∈ {*} ∈ {
        EVENT[at0002] occurrences ∈ {0..1} ∈ {}   -- 1 min sample
        EVENT[at0003] occurrences ∈ {0..1} ∈ {}   -- 2 min sample
        EVENT[at0004] occurrences ∈ {0..1} ∈ {}   -- 3 min sample
    }
}
```

The keyword `cardinality` indicates firstly that the property events must be of a container type, such as `List<T>`, `Set<T>`, `Bag<T>`. The integer range indicates the valid membership of the container; a single '*' means the range 0..*, i.e. '0 to many'. The type of the container is not explicitly indicated, since it is usually defined by the information model. However, the semantics of a logical set (unique membership, ordering not significant), a logical list (ordered, non-unique membership) or a bag (unordered, non-unqique membership) can be constrained using the additional keywords `ordered`, `unordered`, `unique` and `non-unique` within the cardinality constraint, as per the following examples:

```
events cardinality ∈ {*; ordered} ∈ {           -- logical list
events cardinality ∈ {*; unordered; unique} ∈ {  -- logical set
events cardinality ∈ {*; unordered} ∈ {          -- logical bag
```

In theory, none of these constraints can be stronger than the semantics of the corresponding container in the relevant part of the reference model. However, in practice, developers often use lists to facilitate integration, when the actual semantics are intended to be of a set; in such cases, they typically ensure set-like semantics in their own code rather than by using an `Set<T>` type. How such constraints are evaluated in practice may depend somewhat on knowledge of the software system.

> A **cardinality constraint** may be used after any attribute name (or after its existence constraint, if there is one) in order to indicate that the attribute refers to a container type, what number of member items it must have in the data, and optionally, whether it has "list", "set", or "bag" semantics, via the use of the keywords ordered, unordered, unique and non-unique.

Cardinality and existence constraints can co-occur, in order to indicate various combinations on a container type property, e.g. that it is optional, but if present, is a container that may be empty, as in the following:

```
events existence ∈ {0..1} cardinality ∈ {0..*} ∈ {-- etc --}
```

4.3.4.2 Occurrences

A constraint on occurrences is used only with cADL object nodes (not attribute nodes), to indicate how many times in runtime data an instance of a given class conforming to a particular constraint can occur. It only has significance for objects which are children of a container attribute, since by definition, the occurrences of an object which is the value of a single-valued attribute can only be 0..1 or 1..1, and this is already defined by the attribute existence. However, it is not illegal. In the example below, three `EVENT` constraints are shown; the first one ("1 minute sample") is shown as mandatory, while the other two are optional.

```
events cardinality ∈ {*} ∈ {
    EVENT[at0002] occurrences ∈ {1..1} ∈ {}   -- 1 min sample
```

```
            EVENT[at0003] occurrences ∈ {0..1} ∈ {}  -- 2 min sample
            EVENT[at0004] occurrences ∈ {0..1} ∈ {}  -- 3 min sample
        }
```

Another contrived example below expresses a constraint on instances of GROUP such that for GROUPs representing tribes, clubs and families, there can only be one "head", but there may be many members.

```
    GROUP[at0103] ∈ {
        kind ∈ {/tribe|family|club/}
        members cardinality ∈ {*} ∈ {
            PERSON[at0104] occurrences ∈ {1} ∈ {
                title ∈ {"head"}
                -- etc --
            }
            PERSON[at0105] occurrences ∈ {0..*} ∈ {
                title ∈ {"member"}
                -- etc --
            }
        }
    }
```

The first occurrences constraint indicates that a PERSON with the title "head" is mandatory in the GROUP, while the second indicates that at runtime, instances of PERSON with the title "member" can number from none to many. Occurrences may take the value of any range including {0..*}, meaning that any number of instances of the given class may appear in data, each conforming to the one constraint block in the archetype. A single positive integer, or the infinity indicator, may also be used on its own, thus: {2}, {*}. The default occurrences, if none is mentioned, is {1..1}.

> An **occurrences constraint** may appear directly after any type name, in order to indicate how many times data objects conforming to the block introduced by the type name may occur in the data.

Where cardinality constraints are used (remembering that occurrences is always there by default, if not explicitly specified), cardinality and occurrences must always be compatible. The validity rule is:

> **VCOC: cardinality/occurrences validity**: the interval represented by: the sum of all occurrences minimum values - the sum of all occurrences maximum values must be inside the interval of the cardinality.

4.3.5 "Any" Constraints

There are two cases where it is useful to state a completely open, or "any", constraint. The "any" constraint is shown by a single asterisk (*) in braces. The first is when it is desired to show explicitly that some property can have any value, such as in the following:

```
    PERSON[at0001] matches {
        name existence matches {0..1} matches {*}
        -- etc --
    }
```

The "any" constraint on *name* means that any value permitted by the underlying information model is also permitted by the archetype; however, it also provides an opportunity to specify an existence constraint which might be narrower than that in the information model. If the existence constraint is the same, an "any" constraint on a property is equivalent to no constraint being stated at all for that property in the cADL.

The second use of "any" as a constraint value is for types, such as in the following:

```
    ELEMENT[at0004] matches {                -- speed limit
```

```
value matches {
    QUANTITY matches {*}
}
}
}
```

The meaning of this constraint is that in the data at runtime, the value property of ELEMENT must be of type QUANTITY, but can have any value internally. This is most useful for constraining objects to be of a certain type, without further constraining value, and is especially useful where the information model contains subtyping, and there is a need to restrict data to be of certain subtypes in certain contexts.

4.3.6   Object Node Identification and Paths

In many of the examples above, some of the object node typenames are followed by a node identifier, shown in brackets.

> **Node identifiers** are required for any object node which is intended to be addressable elsewhere in the cADL text, or in the runtime system and which would otherwise be ambiguous i.e. has sibling nodes.

In the following example, the PERSON type does not require an identifier, since no sibling node exists at the same level, and unambigous paths can be formed:

```
members cardinality ∈ {*} ∈ {
    PERSON ∈ {
        title ∈ {"head"}
    }
}
```

The path to the *title* attribute is

```
members/title
```

However, where there are more than one sibling node, node identifiers must be used to ensure distinct paths:

```
members cardinality ∈ {*} ∈ {
    PERSON[at0104] ∈ {
        title ∈ {"head"}
    }
    PERSON[at0105] matches {
        title ∈ {"member"}
    }
}
```

The paths to the respective *title* attributes are now:

```
members[at0104]/title
members[at0105]/title
```

Logically, all non-unique parent nodes of an identified node must also be identified back to the root node. The primary function of node identifiers is in forming paths, enabling cADL nodes to be unambiguously referred to. The node identifier can also perform a second function, that of giving a design-time *meaning* to the node, by equating the node identifier to some description. Thus, in the example shown in section 4.3.1, the ELEMENT node is identified by the code [at0010], which can be designated elsewhere in an archetype as meaning "diastolic blood pressure".

Node ids are required only where it is necessary to create paths, for example in "use" statements. However, the underlying reference model might have stronger requirements. The *open*EHR EHR information models [17] for example require that all node types which inherit from the class LOCATA-

BLE have both a *archetype_node_id* and a runtime *name* attribute. Only data types (such as QUANTITY, CODED_TEXT) and their constituent types are exempt.

Paths are used in cADL to refer to cADL nodes, and are expressed in the ADL path syntax, described in detail in section 7 on page 81. ADL paths have the same alternating object/attribute structure implied in the general hierarchical structure of cADL, obeying the pattern TYPE/attribute/TYPE/attribute/…．

> **Paths** in cADL always refer to object nodes, and can only be constructed through nodes having node ids, or nodes which are the only child object of a single-cardinality attribute.

Unusually for a path syntax, a trailing object identifier can be required, even if the attribute corresponds to a single relationship (as might be expected with the "name" property of an object) because in cADL, it is legal to define multiple alternative object constraints - each identified by a unique node id - for a relationship node which has single cardinality.

Consider the following cADL example:

```
HISTORY occurrences ∈ {1} ∈ {
    periodic ∈ {False}
    events cardinality ∈ {*} ∈ {
        EVENT[at0002] occurrences ∈ {0..1} ∈ {}   -- 1 min sample
        EVENT[at0003] occurrences ∈ {0..1} ∈ {}   -- 2 min sample
        EVENT[at0004] occurrences ∈ {0..1} ∈ {}   -- 3 min sample
    }
}
```

The following paths can be constructed:

```
/                       -- the HISTORY object
/periodic               -- the HISTORY.periodic attribute
/events[at0002]         -- the 1 minute event object
/events[at0003]         -- the 2 minute event object
/events[at0004]         -- the 3 minute event object
```

It is valid to add attribute references to the end of a path, if the underlying information model permits it, as in the following example.

```
/events/count           -- count attribute of the items property
```

These examples are *physical* paths because they refer to object nodes using codes. Physical paths can be converted to *logical* paths using descriptive meanings for node identifiers, if defined. Thus, the following two paths might be equivalent:

```
/events[at0004]         -- the 3 minute event object
/events[3 minute event] -- the 3 minute event object
```

None of the paths shown here have any validity outside the cADL block in which they occur, since they do not include an identifier of the enclosing document, normally an archetype. To reference a cADL node in a document from elsewhere (e.g. another archetype of a template) requires that the identifier of the document itself be prefixed to the path, as in the following archetype example:

```
[openehr-ehr-entry.apgar-result.v1]/events[at0002]
```

This kind of path expression is necessary to form the larger paths which occur when archetypes are composed to form larger structures.

4.3.7   Internal References

It occurs reasonably often that one needs to include a constraint which is a repeat of an earlier complex constraint, but within a different block. This is achieved using an archetype internal reference, according to the following rule:

> An **archetype internal reference** is introduced with the `use_node` keyword, in a line of the following form:
>
> `use_node TYPE object_path`

This statement says: use the node of type `TYPE`, found at (the existing) path `object_path`. The following example shows the definitions of the `ADDRESS` nodes for phone, fax and email for a home `CONTACT` being reused for a work `CONTACT`.

```
PERSON ∈ {
    identities ∈ {
        -- etc --
    }
    contacts cardinality ∈ {0..*} ∈ {
        CONTACT [at0002] ∈ {                           -- home address
            purpose ∈ {-- etc --}
            addresses ∈ {-- etc --}
        }
        CONTACT [at0003] ∈ {                           -- postal address
            purpose ∈ {-- etc --}
            addresses ∈ {-- etc --}
        }
        CONTACT [at0004] ∈ {                           -- home contact
            purpose ∈ {-- etc --}
            addresses cardinality ∈ {0..*} ∈ {
                ADDRESS [at0005] ∈ {                   -- phone
                    type ∈ {-- etc --}
                    details ∈ {-- etc --}
                ADDRESS [at0006] ∈ {                   -- fax
                    type ∈ {-- etc --}
                    details ∈ {-- etc --}
                }
                ADDRESS [at0007] ∈ {                   -- email
                    type ∈ {-- etc --}
                    details ∈ {-- etc --}
                }
            }
        }
        CONTACT [at0008] ∈ {                           -- work contact
            purpose ∈ {-- etc --}
            addresses cardinality ∈ {0..*} ∈ {
                use_node ADDRESS /contacts[at0004]/addresses[at0005] -- phone
                use_node ADDRESS /contacts[at0004]/addresses[at0006] -- fax
                use_node ADDRESS /contacts[at0004]/addresses[at0007] -- email
            }
        }
    }
```

4.3.8   Archetype Slots

At any point in a cADL definition, a constraint can be defined which allows other archetypes to be used, rather than defining the desired constraints inline. This is known as an archetype "slot", or

"chaining point", i.e. a connection point whose allowable fillers are constrained by a set of statements, written in the ADL assertion language (defined in section 5 on page 75).

An archetype slot is defined in terms of two lists of assertions statements defining which archetypes are allowed and/or which are excluded from filling that slot. Since archetype slots are typed, the (possibly abstract) type of the allowed archetypes is already constrained. Otherwise, any assertion about a filler archetype can be made. The assertions do not constrain data in the way that other archetype statements do, instead they constrain archetypes. Two kinds of reference may be used in a slot assertion. The first is a reference to an object-oriented property of the filler archetype itself, where the property names are defined by the ARCHETYPE class in the Archetype Object Model. The syntax is typical object-oriented "dot" notation, preceded by an '@' symbol which stands for the archetype. Examples include:

```
@archetype_id
@parent_archetype_id
@short_concept_name
```

This kind of reference is usually used to constrain the *archetype_id* of the filler candidates. The second kind of reference is to absolute paths in the filler archetype, and these take the same form as archetype paths elsewhere in the archetype.

The slot might be "wide", meaning it allows numerous other archetypes, or "narrow", where it allows only a few or just one archetype. All references to attributes of an archetype take the form of absolute paths starting from the root of the archetype. The paths are defined with respect to the archetype object model (AOM). Typical paths include:

A basic kind of assertion is on the identifier of archetypes allowed in the slot, and it is possible to limit this to a single archetype, meaning that the choice of archetype in that slot is fixed. In general, however, the intention of archetypes is to provide highly re-usable models of real world content with local constraining left to templates. The full semantics of archetype chaining are described in the "Archetype Object Model" document.

> An **archetype slot** is introduced with the keyword `allow_archetype`, and is expressed using two lists of assertions, each introduced with the keywords `include` and `exclude`, respectively.

The following example shows how the "Objective" SECTION in a problem/SOAP headings archetype defines two slots, indicating which OBSERVATION and SECTION archetypes are allowed and excluded under the *items* property.

```
SECTION [at2000] occurrences ∈ {0..1} ∈ {          -- objective
    items ∈ {
        allow_archetype OBSERVATION occurrences ∈ {0..1} ∈ {
            include
                @short_concept_name ∈ {/.+/}
        }
        allow_archetype SECTION occurrences ∈ {0..*} ∈ {
            include
                @archetype_id ∈ {/.*\.iso-ehr\.section\..*\..*/}
            exclude
                @archetype_id ∈ {/.*\.iso-ehr\.section\.patient_details\..*/}
        }
    }
}
```

Here, every constraint inside the block starting on an `allow_archetype` line contains constraints that must be met by archetypes in order to fill the slot. (Note that in the examples above, and in all ADL,

the '\' character in paths is not a path separator, as some Microsoft Windows users might assume - it retains its meaning as a quoting character; above it is used to quote the '.' character to ensure it has its literal meaning of '.' rather than "any character" which is its regular expression meaning. Although not particularly beautiful, the syntax used above to indicate allowed values of the *id* attribute is completely standard regular expression syntax, and would be familiar to most users of Unix operating systems (e.g. Linux, BSD etc), Perl, and many other formalisms. Developers of archetypes using GUI tools should of course be spared such technical details.)

Other constraints are possible as well, including that the allowed archetype must contain a certain keyword, or a certain path. The latter is quite powerful – it allows archetypes to be linked together on the basis of context. For example, under a "genetic relatives" heading in a Family History Organiser archetype, the following logical constraint might be used:

```
allow_archetype EVALUATION occurrences ∈ {0..*} matches {
    include
        @short_concept_name ∈ {"family_history_subject"}
          ∧ ∃ /subject/relation
            → /subject/relation ∈ {
               CODED_TEXT ∈ {
                   code ∈ {[ac0003]}              -- "parent" or "sibling"
               }
            }
}
```

4.3.9   Mixed Structures

Three types of structure which represent constraints on complex objects have been presented so far:

- *complex object structures*: any node introduced by a type name and followed by {} containing constraints on attributes;
- *internal references*: any node introduced by the keyword use_node, followed by a type name; such nodes indicate re-use of a complex object constraint that has already been expressed elsewhere in the archetype;
- *archetype slots*: any node introduced by the keyword allow_archetype, followed by a type name; such nodes indicate a complex object constraint which is expressed in some other archetype.

At any given node, all three types can co-exist, as in the following example:

```
SECTION[at2000] ∈ {
    items cardinality ∈ {0..*; ordered} ∈ {
        ENTRY[at2001] ∈ {-- etc --}
        allow_archetype ENTRY ∈ {-- etc --}
        use_node ENTRY [at0001]/some_path[at0004]/
        ENTRY[at2002] ∈ {-- etc --}
        use_node ENTRY /[at1002]/some_path[at1012]/
        use_node ENTRY /[at1005]/some_path[at1052]/
        ENTRY[at2003] ∈ {-- etc --}
    }
}
```

Here, we have a constraint on an attribute called *items* (of cardinality 0..*), expressed as a series of possible constraints on objects of type ENTRY. The 1st, 4th and 7th are described "in place" (the details are removed here, for brevity); the 3rd, 5th and 6th are expressed in terms of internal references to other nodes earlier in the archetype, while the 2nd is an archetype slot, whose constraints are expressed in other archetypes matching the include/exclude constraints appearing between the braces

of this node (again, avoided for the sake of brevity). Note also that the `ordered` keyword has been used to indicate that the list order is intended to be significant.

4.4 Constraints on Primitive Types

While constraints on complex types follow the rules described so far, constraints on attributes of primitive types in cADL can be expressed without type names, and omitting one level of braces, as follows:

```
some_attr matches {some_pattern}
```

rather than:

```
some_attr matches {
    BASIC_TYPE matches {
        some_pattern
    }
}
```

This is made possible because the syntax patterns of all primitive types constraints are mutually distinguishable, i.e. the type can always be inferred from the syntax alone. Since all leaf attributes of all object models are of primitive types, or lists or sets of them, cADL archetypes using the brief form for primitive types are significantly less verbose overall, as well as being more directly comprehensible to human readers. cADL does not however oblige the brief form described here, and the more verbose one can be used. In either case, the syntax of the pattern appearing within the final pair of braces obeys the rules described below.

4.4.1 Constraints on String

Strings can be constrained in two ways: using a list of fixed strings, and using using a regular expression. All constraints on strings are case-sensitive.

4.4.1.1 List of Strings

A String-valued attribute can be constrained by a list of strings (using the dADL syntax for string lists), including the simple case of a single string. Examples are as follows:

```
species ∈ {"platypus"}
species ∈ {"platypus", "kangaroo"}
species ∈ {"platypus", "kangaroo", "wombat"}
```

The first example constraints the runtime value of the *species* attribute of some object to take the value "platypus"; the second constrains it be either "platypus" or "kangaroo", and so on. **In almost all cases, this kind of string constraint should be avoided**, since it usually renders the body of the archetype language-dependent. Exceptions are proper names (e.g. "NHS", "Apgar"), product trade-names (but note even these are typically different in different language locales, even if the different names are not literally translations of each other). The preferred way of constraining string attributes in a language independent way is with local [ac] codes. See Local Constraint Codes on page 86.

4.4.1.2 Regular Expression

The second way of constraining strings is with regular expressions, a widely used syntax for expressing patterns for matching strings. The regular expression syntax used in cADL is a proper subset of that used in the Perl language (see [18] for a full specification of the regular expression language of Perl). Three uses of it are accepted in cADL:

```
string_attr matches {/regular expression/}
string_attr matches {=~ /regular expression/}
string_attr matches {!~ /regular expression/}
```

The first two are identical, indicating that the attribute value must match the supplied regular expression. The last indicates that the value must *not* match the expression. If the delimiter character is required in the pattern, it must be quoted with the backslash ('\') character, or else alternative delimiters can be used, enabling more comprehensible patterns. A typical example is regular expressions including units. The following two patterns are equivalent:

```
units ∈ {/km\/h|mi\/h/}
units ∈ {^km/h|mi/h^}
```

The rules for including special characters within strings follow those for dADL. In regular expressions, the small number of special characters are quoted according to the rules of Perl regular expressions; all other characters are quoted using the ISO and XML conventions described in the section on dADL.

The regular expression patterns supported in cADL are as follows.

Atomic Items

.   match any single character. E.g. / `...` / matches any 3 characters which occur with a space before and after;

[xyz]   match any of the characters in the set `xyz` (case sensitive). E.g. `/[0-9]/` matches any string containing a single decimal digit;

[a-m]   match any of the characters in the set of characters formed by the continuous range from `a` to `m` (case sensitive). E.g. `/[0-9]/` matches any single character string containing a single decimal digit, /[S-Z]/ matches any single character in the range `S` - `z`;

[^a-m]   match any character except those in the set of characters formed by the continuous range from `a` to `m`. E.g. `/[^0-9]/` matches any single character string as long as it does not contain a single decimal digit;

Grouping

(pattern) parentheses are used to group items; any pattern appearing within parentheses is treated as an atomic item for the purposes of the occurrences operators. E.g. `/([0-9][0-9])/` matches any 2-digit number.

Occurrences

*   match 0 or more of the preceding atomic item. E.g. `/.*/` matches any string; `/[a-z]*/` matches any non-empty lower-case alphabetic string;

+   match 1 or more occurrences of the preceding atomic item. E.g. `/a.+/` matches any string starting with 'a', followed by at least one further character;

?   match 0 or 1 occurrences of the preceding atomic item. E.g. `/ab?/` matches the strings "a" and "ab";

{m,n}   match m to n occurrences of the preceding atomic item. E.g. `/ab{1,3}/` matches the strings "ab" and "abb" and "abbb"; `/[a-z]{1,3}/` matches all lower-case alphabetic strings of one to three characters in length;

{m,}   match at least m occurrences of the preceding atomic item;

{,n}   match at most n occurrences of the preceding atomic item;

{m}   match exactly m occurrences of the preceding atomic item;

Special Character Classes

\d, \D   match a decimal digit character; match a non-digit character;

    `\s`, `\S` match a whitespace character; match a non-whitespace character;

Alternatives

    `pattern1|pattern2`   match either pattern1 or pattern2. E.g. `/lying|sitting|standing/` matches any of the words "`lying`", "`sitting`" and "`standing`".

A similar warning should be noted for the use of regular expressions to constrain strings: they should be limited to non-linguistically dependent patterns, such as proper and scientific names. The use of regular expressions for constraints on normal words will render an archetype linguistically dependent, and potentially unusable by others.

4.4.2 Constraints on Integer

Integers can be constrained using a list of integer values, and using an integer interval.

4.4.2.1 List of Integers

Lists of integers expressed in the syntax from dADL (described in Lists of Built-in Types on page 34) can be used as a constraint, e.g.:

```
length matches {1000}          -- fixed value of 1000
magnitude matches {0, 5, 8}    -- any of 0, 5 or 8
```

The first constraint requires the attribute length to be 1000, while the second limits the value of magnitude to be 0, 5, or 8 only.

4.4.2.2 Interval of Integer

Integer intervals are expressed using the interval syntax from dADL (described in Intervals of Ordered Primitive Types on page 33). Examples include

```
length matches {|1000|}        -- point interval of 1000 (=fixed value)
length matches {|950..1050|}   -- allow 950 - 1050
length matches {|0..1000|}     -- allow 0 - 1000
length matches {|<10|}         -- allow up to 9
length matches {|>10|}         -- allow 11 or more
length matches {|<=10|}        -- allow up to 10
length matches {|>=10|}        -- allow 10 or more
length matches {|100+/-5|}     -- allow 100 +/- 5, i.e. 95 - 105
rate matches {|0..infinity|}   -- allow 0 - infinity, i.e. same as >= 0
```

Intervals may be combined in integer constraints, using the semicolon character (';') as follows:

```
critical_range matches {|5..9; 101..110|}
```

4.4.3 Constraints on Real

Constraints on Real values follow exactly the same syntax as for Integers, in both list and interval forms. The only difference is that the real number values used in the constraints are indicated by the use of the decimal point and at least one succeeding digit, which may be 0. Typical examples are:

```
magnitude ∈ {5.5}                    -- fixed value
magnitude ∈ {|5.5|}                  -- point interval (=fixed value)
magnitude ∈ {|5.5..6.0|}             -- interval
magnitude ∈ {5.5, 6.0, 6.5}          -- list
magnitude ∈ {|<10.0|}                -- allow anything less than 10.0
magnitude ∈ {|>10.0|}                -- allow greater than 10.0
magnitude ∈ {|<=10.0|}               -- allow up to 10.0
magnitude ∈ {|>=10.0|}               -- allow 10.0 or more
magnitude ∈ {|80.0+/-12.0|}          -- allow 80 +/- 12
```

4.4.4   Constraints on Boolean

Boolean runtime values can be constrained to be True, False, or either, as follows:

```
some_flag matches {True}
some_flag matches {False}
some_flag matches {True, False}
```

4.4.5   Constraints on Character

Characters can be constrained in two ways: using a list of characters, and using a regular expression.

4.4.5.1   List of Characters

The following examples show how a character value may be constrained using a list of fixed character values. Each character is enclosed in single quotes.

```
color_name matches {'r'}
color_name matches {'r', 'g', 'b'}
```

4.4.5.2   Regular Expression

Character values can also be constrained using single-character regular expression elements, also enclosed in single quotes, as per the following examples:

```
color_name matches {'[rgbcmyk]'}
color_name matches {'[^\s\t\n]'}
```

The only allowed elements of the regular expression syntax in character expressions are the following:

- any item from the Atomic Items list above;
- any item from the Special Character Classes list above;
- the '.' character, standing for "any character";
- an alternative expression whose parts are any item types, e.g. `'a'|'b'|[m-z]`

4.4.6   Constraints on Dates, Times and Durations

Dates, times, date/times and durations may all be constrained in three ways: using a list of values, using intervals, and using patterns. The first two ways allow values to be constrained to actual date, time etc values, while the last allows values to be constrained on the basis of which parts of the date, time etc are present or missing, regardless of value. The pattern method is described first, since patterns can also be used in lists and intervals.

> **NB:** for the following constraint types, parser writers should consider allowing the 'T' character to be optional on read but mandatory on save for some time. This is because previous versions of ADL did not include it, with the result that existing tools have created archetypes without the 'T' in date/time constraint patterns.

4.4.6.1   Date, Time and Date/Time

Patterns

Dates, times, and date/times (i.e. timestamps), can be constrained using patterns based on the ISO 8601 date/time syntax, which indicate which parts of the date or time must be supplied. A constraint pattern is formed from the abstract pattern `yyyy-mm-ddThh:mm:ss` (itself formed by translating each field of an ISO 8601 date/time into a letter representing its type), with either '?' (meaning optional) or 'X' (not allowed) characters substituted in appropriate places. A simplified grammar of the pattern is as follows (EBNF; all tokens shown are literals):

```
date_constraint:        yyyy - mm|??|XX - dd|??|XX
```

```
time_constraint:          T hh : mm|??|XX : ss|??|XX
time_in_date_constraint:  T hh|??|XX : mm|??|XX : ss|??|XX
date_time_constraint:     date_constraint time_in_date_constraint
```

All expressions generated by this grammar must also satisfy the validity rules:

- where '??' appears in a field, only '??' or 'XX' can appear in fields to the right
- where 'XX' appears in a field, only 'XX' can appear in fields to the right

A fuller grammar can be defined to implement both the simplifed grammar and validity rules.

The following table shows the valid patterns that can be used, and the types implied by each pattern.

Implied Type	Pattern	Explanation
Date	yyyy-mm-dd	full date must be specified
Date, Partial Date	yyyy-mm-??	optional day; e.g. day in month forgotten
Date, Partial Date	yyyy-??-??	optional month, day; i.e. any date allowed; e.g. mental health questionnaires which include well known historical dates
Partial Date	yyyy-??-XX	optional month, no day; (any examples?)
Time	Thh:mm:ss	full time must be specified
Partial Time	Thh:mm:XX	no seconds; e.g. appointment time
Partial Time	Thh:??:XX	optional minutes, no seconds; e.g. normal clock times
Time, Partial Time	Thh:??:??	optional minutes, seconds; i.e. any time allowed
Date/Time	yyyy-mm-ddThh:mm:ss	full date/time must be specified
Date/Time, Partial Date/Time	yyyy-mm-ddThh:mm:??	optional seconds; e.g. appointment date/time
Partial Date/Time	yyyy-mm-ddThh:mm:XX	no seconds; e.g. appointment date/time
Partial Date/Time	yyyy-mm-ddThh:??:XX	no seconds, minutes optional; e.g. in patient-recollected date/times
Date/Time, Partial Date/Time Partial Date/Partial Time	yyyy-??-??T??:??:??	minimum valid date/time constraint

List and Intervals

Dates, times and date/times can also be constrained using lists and intervals. Each date, time etc in such a list or interval may be a literal date, time etc value, or a value based on a pattern. In the latter case, the limit values are specified using the patterns from the above table, but with numbers in the positions where 'x' and '?' do not appear. For example, the pattern yyyy-??-XX could be transformed into 1995-??-XX to mean any partial date in 1995. Examples of such constraints:

```
|1995-??-XX|             -- any partial date in 1995
|T09:30:00|              -- exactly 9:30 am
|< T09:30:00|            -- any time before 9:30 am
```

```
|<= T09:30:00|                              -- any time at or before 9:30 am
|> T09:30:00|                               -- any time after 9:30 am
|>= T09:30:00|                              -- any time at or after 9:30 am
|2004-05-20..2004-06-02|                    -- a date range
|2004-05-20T00:00:00..2005-05-19T23:59:59| -- a date/time range
```

4.4.6.2   Duration Constraints

Patterns

Patterns based on ISO 8601 can be used to constraint duratoins in the same way as for Date/time types. The general form of a pattern is (EBNF; all tokens are literals):

```
P[Y|y][M|m][W|w][D|d][T[H|h][M|m][S|s]]
```

Note that allowing the 'W' designator to be used with the other designators corresponds to a deviation from the published ISO 8601 standard used in *open*EHR, namely:

- durations are supposed to take the form of PnnW or PnnYnnMnnDTnnHnnMnnS, but in *open*EHR, the W (week) designator can be used with the other designators, since it is very common to state durations of pregnancy as some combination of weeks and days.

The use of this pattern indicates which "slots" in an ISO duration string may be filled. Where multiple letters are supplied in a given pattern, the meaning is "or", i.e. any one or more of the slots may be supplied in the data. This syntax allows specifications like the following to be made:

```
Pd          -- a duration containing days only, e.g. P5d
Pm          -- a duration containing months only, e.g. P5m
PTm         -- a duration containing minutes only, e.g. PT5m
Pwd         -- a duration containing weeks and/or days only, e.g. P4w
PThm        -- a duration containing hours and/or minutes only, e.g. PT2h30m
```

List and Intervals

Durations can also be constrained by using absolute ISO 8601 duration values, or ranges of the same, e.g.:

```
PT1m              -- 1 minute
P1dT8h            -- 1 day 8 hrs
|PT0m..PT1m30s|   -- Reasonable time offset of first apgar sample
```

4.4.7   Constraints on Lists of Primitive types

In many cases, the type in the information model of an attribute to be constrained is a list or set of primitive types, e.g. List<Integer>, Set<String> etc. As for complex types, this is indicated in cADL using the `cardinality` keyword (as for complex types), as follows:

```
some_attr cardinality ∈ {0..*} ∈ {some_constraint}
```

The pattern to match in the final braces will then have the meaning of a list or set of value constraints, rather than a single value constraint. Any constraint described above for single-valued attributes, which is commensurate with the type of the attribute in question, may be used. However, as with complex objects, the meaning is now that every item in the list is constrained to be any one of the values implied by the constraint expression. For example,

```
speed_limits cardinality ∈ {0..*; ordered} ∈ {50, 60, 70, 80, 100, 130}
```

constrains each value in the list corresponding to the value of the attribute *speed_limits* (of type List<Integer>), to be any one of the values 50, 60, 70 etc.

4.4.8 Assumed Values

When archetypes are defined to have optional parts, an ability to define 'assumed' values is useful. For example, an archetype for the concept 'blood pressure measurement' might contain an optional protocol section describing the patient position, with choices 'lying', 'sitting' and 'standing'. Since the section is optional, data could be created according to the archetype which does not contain the protocol section. However, a blood pressure cannot be taken without the patient in some position, so clearly there could be an implied or 'assumed' value.

The archetype allows this to be explicitly stated so that all users/systems know what value to assume when optional items are not included in the data. Assumed values are optionally definable on primitive types only, and are expressed after the constraint expression, by a semi-colon (';') followed by a value of the same type as that implied by the preceding part of the constraint. The use of assumed values is illustrated here for a number of primitive types:

```
length matches {|0..1000|; 200}           -- allow 0 - 1000, assume 200
some_flag matches {True, False; True}      -- allow T or F, assume T
some_date matches {yyyy-mm-ddThh:mm:XX; 1800-01-01T00:00:00}
```

If no assumed value is stated, no reliable assumption can be made by the receiver of the archetyped data about what the values of removed optional parts might be, from inspecting the archetype. However, this usually corresponds to a situation where the assumed value does not even need to be stated - the same value will be assumed by all users of this data, if its value is not transmitted. In other cases, it may be that it doesn't matter what the assumed value is. For example, an archetype used to capture physical measurements might include a "protocol" section, which in turn can be used to record the "instrument" used to make a given measurement. In a blood pressure specialisation of this archetype it is fairly likely that physicians recording or receiving the data will not care about what instrument was used.

4.5 Syntax Specification

The cADL grammar is available as an HTML document. This grammar is implemented and tested using lex (.l file) and yacc (.y file) specifications for in the Eiffel programming environment. The current release of these files is available at http://my.openehr.org/wsvn/ref_impl_eiffel/TRUNK/components/adl_parser/src/syntax/cadl/parser/?rev=0&sc=0. The .l and .y files can easily be converted for use in another yacc/lex-based programming environment.

4.5.1 Grammar

The following provides the cADL parser production rules (yacc specification) as of revision 36 of the Eiffel reference implementation repository (http://svn.openehr.org/ref_impl_eiffel). Note that because of interdependcies with path and assertion production rules, practical implementations may have to include all production rules in one parser.

```
input:
  c_complex_object
| error


c_complex_object:
  c_complex_object_head SYM_MATCHES SYM_START_CBLOCK c_complex_object_body
SYM_END_CBLOCK


c_complex_object_head:
  c_complex_object_id c_occurrences
```

c_complex_object_id:
  V_TYPE_IDENTIFIER
| V_TYPE_IDENTIFIER V_LOCAL_TERM_CODE_REF

c_complex_object_body:
  c_any
| c_attributes

c_object:
  c_complex_object
| archetype_internal_ref
| archetype_slot
| constraint_ref
| c_coded_term
| c_ordinal
| c_primitive_object
| V_C_DOMAIN_TYPE
| ERR_C_DOMAIN_TYPE
| error

archetype_internal_ref:
  SYM_USE_NODE V_TYPE_IDENTIFIER object_path
| SYM_USE_NODE V_TYPE_IDENTIFIER error

archetype_slot:
  c_archetype_slot_head SYM_MATCHES SYM_START_CBLOCK c_includes c_excludes
SYM_END_CBLOCK

c_archetype_slot_head:
  c_archetype_slot_id c_occurrences

c_archetype_slot_id:
  SYM_ALLOW_ARCHETYPE V_TYPE_IDENTIFIER
| SYM_ALLOW_ARCHETYPE V_TYPE_IDENTIFIER V_LOCAL_TERM_CODE_REF
| SYM_ALLOW_ARCHETYPE error

c_primitive_object:
  c_primitive

c_primitive:
  c_integer
| c_real
| c_date
| c_time
| c_date_time
| c_duration
| c_string
| c_boolean
| error

c_any:
  *

c_attributes:
  c_attribute
| c_attributes c_attribute

```
    c_attribute:
      c_attr_head SYM_MATCHES SYM_START_CBLOCK c_attr_values SYM_END_CBLOCK

    c_attr_head:
      V_ATTRIBUTE_IDENTIFIER c_existence
    | V_ATTRIBUTE_IDENTIFIER c_existence c_cardinality

    c_attr_values:
      c_object
    | c_attr_values c_object
    | c_any
    | error

    c_includes:
      -/-
    | SYM_INCLUDE assertions

    c_excludes:
      -/-
    | SYM_EXCLUDE assertions

    c_existence:
      -/-
    | SYM_EXISTENCE SYM_MATCHES SYM_START_CBLOCK existence_spec SYM_END_CBLOCK

    existence_spec:
      V_INTEGER
    | V_INTEGER SYM_ELLIPSIS V_INTEGER

    c_cardinality:
      SYM_CARDINALITY SYM_MATCHES SYM_START_CBLOCK cardinality_spec
    SYM_END_CBLOCK

    cardinality_spec:
      occurrence_spec
    | occurrence_spec ; SYM_ORDERED
    | occurrence_spec ; SYM_UNORDERED
    | occurrence_spec ; SYM_UNIQUE
    | occurrence_spec ; SYM_ORDERED ; SYM_UNIQUE
    | occurrence_spec ; SYM_UNORDERED ; SYM_UNIQUE
    | occurrence_spec ; SYM_UNIQUE ; SYM_ORDERED
    | occurrence_spec ; SYM_UNIQUE ; SYM_UNORDERED

    cardinality_limit_value:
      integer_value
    | *

    c_occurrences:
      -/-
    | SYM_OCCURRENCES SYM_MATCHES SYM_START_CBLOCK occurrence_spec
    SYM_END_CBLOCK
    | SYM_OCCURRENCES error

    occurrence_spec:
      cardinality_limit_value
    | V_INTEGER SYM_ELLIPSIS cardinality_limit_value
```

```
c_integer_spec:
  integer_value
| integer_list_value
| integer_interval_value
| occurrence_spec

c_integer:
  c_integer_spec
| c_integer_spec ; integer_value
| c_integer_spec ; error

c_real_spec:
  real_value
| real_list_value
| real_interval_value

c_real:
  c_real_spec
| c_real_spec ; real_value
| c_real_spec ; error

c_date_spec:
  V_ISO8601_DATE_CONSTRAINT_PATTERN
| date_value
| date_interval_value

c_date:
  c_date_spec
| c_date_spec ; date_value
| c_date_spec ; error

c_time_spec:
  V_ISO8601_TIME_CONSTRAINT_PATTERN
| time_value
| time_interval_value

c_time:
  c_time_spec
| c_time_spec ; time_value
| c_time_spec ; error

c_date_time_spec:
  V_ISO8601_DATE_TIME_CONSTRAINT_PATTERN
| date_time_value
| date_time_interval_value

c_date_time:
  c_date_time_spec
| c_date_time_spec ; date_time_value
| c_date_time_spec ; error

c_duration_spec:
  V_ISO8601_DURATION_CONSTRAINT_PATTERN
| duration_value
| duration_interval_value

c_duration:
```

```
  c_duration_spec
| c_duration_spec ; duration_value
| c_duration_spec ; error
```

c_string_spec:
```
  V_STRING
| string_list_value
| string_list_value , SYM_LIST_CONTINUE
| V_REGEXP
```

c_string:
```
  c_string_spec
| c_string_spec ; string_value
| c_string_spec ; error
```

c_boolean_spec:
```
  SYM_TRUE
| SYM_FALSE
| SYM_TRUE , SYM_FALSE
| SYM_FALSE , SYM_TRUE
```

c_boolean:
```
  c_boolean_spec
| c_boolean_spec ; boolean_value
| c_boolean_spec ; error
```

c_ordinal:
```
  c_ordinal_spec
| c_ordinal_spec ; integer_value
| c_ordinal_spec ; error
```

c_ordinal_spec:
```
  ordinal
| c_ordinal_spec , ordinal
```

ordinal:
```
  integer_value SYM_INTERVAL_DELIM V_QUALIFIED_TERM_CODE_REF
```

c_coded_term:
```
  V_TERM_CODE_CONSTRAINT
| V_QUALIFIED_TERM_CODE_REF
```

constraint_ref:
```
  V_LOCAL_TERM_CODE_REF
```

any_identifier:
```
  V_TYPE_IDENTIFIER
| V_ATTRIBUTE_IDENTIFIER
```

```
-- for string_value etc, see dADL spec

-- for attribute_path, object_path, call_path, etc, see Path spec

-- for assertions, assertion, see Assertion spec
```

4.5.2   Symbols

The following shows the lexical specification for the cADL grammar.

```
    ----------/* comments */ -----------------------------------------
    "--".*                                        -- Ignore comments
    "--".*\n[ \t\r]*

    ----------/* symbols */ ------------------------------------------
    "-"          -- -> Minus_code
    "+"          -- -> Plus_code
    "*"          -- -> Star_code
    "/"          -- -> Slash_code
    "^"          -- -> Caret_code
    "="          -- -> Equal_code
    "."          -- -> Dot_code
    ";"          -- -> Semicolon_code
    ","          -- -> Comma_code
    ":"          -- -> Colon_code
    "!"          -- -> Exclamation_code
    "("          -- -> Left_parenthesis_code
    ")"          -- -> Right_parenthesis_code
    "$"          -- -> Dollar_code

    "??"         -- -> SYM_DT_UNKNOWN
    "?"          -- -> Question_mark_code

    "|"          -- -> SYM_INTERVAL_DELIM

    "["          -- -> Left_bracket_code
    "]"          -- -> Right_bracket_code

    "{"          -- -> SYM_START_CBLOCK
    "}"          -- -> SYM_END_CBLOCK

    ".."         -- -> SYM_ELLIPSIS
    "..."        -- -> SYM_LIST_CONTINUE

    ----------/* common keywords */ -----------------------------------

    [Mm][Aa][Tt][Cc][Hh][Ee][Ss] -- -> SYM_MATCHES

    [Ii][Ss]_[Ii][Nn]            -- -> SYM_MATCHES

    ----------/* assertion keywords */ --------------------------------

    [Tt][Hh][Ee][Nn]             -- -> SYM_THEN

    [Ee][Ll][Ss][Ee]             -- -> SYM_ELSE

    [Aa][Nn][Dd]                 -- -> SYM_AND

    [Oo][Rr]                     -- -> SYM_OR

    [Xx][Oo][Rr]                 -- -> SYM_XOR

    [Nn][Oo][Tt]                 -- -> SYM_NOT

    [Ii][Mm][Pp][Ll][Ii][Ee][Ss]  -- -> SYM_IMPLIES

    [Tt][Rr][Uu][Ee]             -- -> SYM_TRUE
```

```
[Ff][Aa][Ll][Ss][Ee]                -- -> SYM_FALSE

[Ff][Oo][Rr][_][Aa][Ll][Ll]     -- -> SYM_FORALL

[Ee][Xx][Ii][Ss][Tt][Ss]          -- -> SYM_EXISTS

---------/* cADL keywords */ -----------------------------------------

[Ee][Xx][Ii][Ss][Tt][Ee][Nn][Cc][Ee]              -- -> SYM_EXISTENCE

[Oo][Cc][Cc][Uu][Rr][Rr][Ee][Nn][Cc][Ee][Ss]      -- -> SYM_OCCURRENCES

[Cc][Aa][Rr][Dd][Ii][Nn][Aa][Ll][Ii][Tt][Yy]      -- -> SYM_CARDINALITY

[Oo][Rr][Dd][Ee][Rr][Ee][Dd]                       -- -> SYM_ORDERED

[Uu][Nn][Oo][Rr][Dd][Ee][Rr][Ee][Dd]              -- -> SYM_UNORDERED

[Uu][Nn][Ii][Qq][Uu][Ee]                           -- -> SYM_UNIQUE

[Ii][Nn][Ff][Ii][Nn][Ii][Tt][Yy]                  -- -> SYM_INFINITY

[Uu][Ss][Ee][_][Nn][Oo][Dd][Ee]                   -- -> SYM_USE_NODE

[Uu][Ss][Ee][_][Aa][Rr][Cc][Hh][Ee][Tt][Yy][Pp][Ee] -- ->
SYM_ALLOW_ARCHETYPE

[Aa][Ll][Ll][Oo][Ww][_][Aa][Rr][Cc][Hh][Ee][Tt][Yy][Pp][Ee]
SYM_ALLOW_ARCHETYPE

[Ii][Nn][Cc][Ll][Uu][Dd][Ee]                       -- -> SYM_INCLUDE

[Ee][Xx][Cc][Ll][Uu][Dd][Ee]                       -- -> SYM_EXCLUDE

----------/* V_URI */ -------------------------------------------------
[a-z]+:\/\/[^> ]*

---------/* V_QUALIFIED_TERM_CODE_REF */ ----------------------------
-- any qualified code, e.g. [local::at0001], [local::ac0001], [loinc::700-0]
--

\[[a-zA-Z0-9()._\-]+::[a-zA-Z0-9._\-]+\]

---------/* V_TERM_CODE_CONSTRAINT of form */ ----------------------
-- [terminology_id::code, -- comment
--          code, -- comment
--          code] -- comment
--
-- Form with assumed value
-- [terminology_id::code, -- comment
--          code; -- comment
--          code] -- an optional assumed value
--

\[[a-zA-Z0-9()._\-]+::[ \t\n]*                          -- pick up [ line

<IN_TERM_CONSTRAINT>[ \t]*[a-zA-Z0-9._\-]+[ \t]*;[ \t\n]* -- pick up , line
```

```
<IN_TERM_CONSTRAINT>[ \t]*[a-zA-Z0-9._\-]+[ \t]*,[ \t\n]*    -- pick up ; line

<IN_TERM_CONSTRAINT>\-\-[^\],\n]*\n                          -- do nothing

<IN_TERM_CONSTRAINT>[ \t]*[a-zA-Z0-9._\-]*[ \t\n]*\]   -- pick up ] line


---------/* V_LOCAL_TERM_CODE_REF */ --------------------------------
-- any unqualified code, e.g. [at0001], [ac0001], [700-0]
--
\[[a-zA-Z0-9][a-zA-Z0-9._\-]*\]

---------/* V_LOCAL_CODE */ -----------------------------------------
a[ct][0-9.]+

---------/* V_QUALIFIED_TERM_CODE_REF */ --------------------
-- any qualified code, e.g. [local::at0001], [local::ac0001], [loinc::700-0]
--
\[[a-zA-Z0-9()._\-]+::[a-zA-Z0-9._\-]+\]

---------/* V_ISO8601_DURATION */ ------------------------
P[0-9]+[wW]|
P([0-9]+[yY])?([0-9]+[mM])?([0-9]+[dD])?T([0-9]+[hH])?([0-9]+[mM])?([0-9]+[sS])? |
P([0-9]+[yY])?([0-9]+[mM])?([0-9]+[dD])?

---------/* V_ISO8601_DATE_CONSTRAINT_PATTERN */ ----------------
[yY][yY][yY][yY]-[mM?X][mM?X]-[dD?X][dD?X]

---------/* V_ISO8601_TIME_CONSTRAINT_PATTERN */ ----------------
T[hH][hH]:[mM?X][mM?X]:[sS?X][sS?X]

---------/* V_ISO8601_DATE_TIME_CONSTRAINT_PATTERN */ -------------
[yY][yY][yY][yY]-[mM?][mM?]-[dD?X][dD?X][ T][hH?X][hH?X]:[mM?X][mM?X]:[sS?X][sS?X]

---------/* V_ISO8601_DURATION_CONSTRAINT_PATTERN */ -------------
P[wW]|
P[yY]?[mM]?[dD]?T[hH]?[mM]?[sS]? |
P[yY]?[mM]?[dD]?

---------/* V_TYPE_IDENTIFIER */ ---------------------------------------
[A-Z][a-zA-Z0-9_]*

---------/* V_ATTRIBUTE_IDENTIFIER */ --------------------------------
[a-z][a-zA-Z0-9_]*

---------/* V_C_DOMAIN_TYPE - sections of dADL syntax */ ----------------
{mini-parser specification}

-- this is an attempt to match a dADL section inside cADL. It will
-- probably never work 100% properly since there can be '>' inside "||"
-- ranges, and also strings containing any character, e.g. units string
-- contining "{}" chars. The real solution is to use the dADL parser on
-- the buffer from the current point on and be able to fast-forward the
-- cursor to the last character matched by the dADL scanner

[A-Z][a-zA-Z0-9_]*[ \n]*<                      -- match a pattern like
```

```
                                               -- 'Type_Identifier whitespace <'

    <IN_C_DOMAIN_TYPE>[^}>]*>[ \n]*[^>}A-Z]     -- match up to next > not
                                               -- followed by a '}' or '>'

    <IN_C_DOMAIN_TYPE>[^}>]*>+[ \n]*[}A-Z]      -- final section - '...>
                                               -- whitespace } or beginning of
                                               -- a type identifier'

    <IN_C_DOMAIN_TYPE>[^}>]*[ \n]*}             -- match up to next '}' not
                                               -- preceded by a '>'


    ----------/* V_REGEXP */ -------------------------------------
    {mini-parser specification}
    "{/"                            -- start of regexp
    <IN_REGEXP1>[^/]*\\\/           -- match any segments with quoted slashes
    <IN_REGEXP1>[^/}]*\/            -- match final segment

    \^[^^\n]*\^{                    -- regexp formed using '^' delimiters


    ----------/* V_INTEGER */ -------------------------------------------
    [0-9]+


    ----------/* V_REAL */ ----------------------------------------------
    [0-9]+\.[0-9]+
    [0-9]+\.[0-9]+[eE][+-]?[0-9]+


    ----------/* V_STRING */ --------------------------------------------
    \"[^\\\n"]*\"

    -- strings containing quotes, special characters etc
    {mini-parser specification}
    \"[^\\\n"]*                     -- beginning of a string
    <IN_STR>\\\\                     -- \\ - append '\'
    <IN_STR>\\\"                     -- \" - append '"'
    <IN_STR>&[a-zA-Z][a-zA-Z0-9_]*; -- match ISO special character
                                    -- pattern &char_name;
    <IN_STR>&#x([a-fA-F0-9_]){4};   -- match W3C XML special character
                                    -- pattern &#xHHHH;
    <IN_STR>[^\\\n"]+
    <IN_STR>\\\n[ \t\r]*            -- match LF in line
    <IN_STR>[^\n"]*\"               -- match final end of string

    <IN_STR>.|\n|                   -- Error
```

5   Assertions

5.1   Overview

This section describes the an initial version of the assertion sub-language of archetypes. Assertions are used in archetype "slot" clauses in the cADL `definition` section, and in the `invariant` section. The following simple assertion in an invariant clause says that the speed in kilometres of some node is related to the speed-in-miles by a factor of 1.6:

```
invariant = <
    validity: /speed[at0002]/kilometres/magnitude =
              /speed[at0004]/miles/magnitude * 1.6
>
```

The archetype assertion language is a small language of its own. Formally it is a first-order predicate logic with equality and comparison operators (=, >, etc). It is very nearly a subset of the OMG's emerging OCL (Object Constraint Language) syntax, and is very similar to the assertion syntax which has been used in the Object-Z [14] and Eiffel [12] languages and tools for over a decade. (See Sowa [15], Hein [8], Kilov & Ross [9] for an explanation of predicate logic in information modelling). Further work will be done to more completely define the assertion language.

5.2   Keywords

The syntax of the invariant section is a subset of first-order predicate logic. In it, the following keywords can be used:

- `exists`, `for_all`,
- `and`, `or`, `xor`, `not`, `implies`
- `true`, `false`

Symbol equivalents for some of the above are given in the following table.

Textual Rendering	Symbolic Rendering	Meaning
matches, is_in	$\in$	Set membership, "p is in P"
exists	$\exists$	Existential quantifier, "there exists ..."
for_all	$\forall$	Universal quantifier, "for all x..."
implies	$\supset, \rightarrow$	Material implication, "p implies q", or "if p then q"
and	$\wedge$	Logical conjunction, "p and q"
or	$\vee$	Logical disjunction, "p or q"
xor	$\underline{\vee}$	Exclusive or, "only one of p or q"
not, ~	$\sim, \neg$	Negation, "not p"

The not operator can be applied as a prefix operator to all other operators except `for_all`; either textual rendering "not" or "~" can be used.

5.3   Operators

Assertion expressions can include arithmetic, relational and boolean operators, plus the existential and universal quantifiers.

5.3.1　Arithmetic Operators

The supported arithmetic operators are as follows:

addition: +

subtraction: -

multiplication: *

division: /

exponent: ^

modulo division: % - remainder after integer division

5.3.2　Equality Operators

The supported equality operators are as follows:

equality: =

inequality: <>

The semantics of these operators are of value comparison.

5.3.3　Relational Operators

The supported relational operators are as follows:

less than: <

less than or equal: <=

greater than: >

greater than or equal: >=

The semantics of these operators are of value comparison. Their domain is limited to values of comparable types.

5.3.4　Boolean Operators

The supported boolean operators are as follows:

not: **not**

and: **and**

xor: **xor**

implies: **implies**

set membership: **matches**, **is_in**

The boolean operators also have symbolic equivalents shown earlier.

5.3.5　Quantifiers

The two standard logical quantifier operators are supported:

existential quantifier: **exists**

universal quantifier: **for_all**

These operators also have the usual symbolic equivalents shown earlier.

5.4 Operands

Operands in an assertion expression can be any of the following:

manifest constant: any constant of any primitive type, expressed according to the dADL syntax for values

variable reference: any name starting with '$', e.g. $body_weight;

object reference: a path referring to an object node, i.e. any path ending in a node identifier

property reference: a path referring to a property, i.e. any path ending in ".property_name"

If an assertion is used in an archetype slot definition, its paths refer to the archetype filling the slot, not the one containing the slot.

5.5 Precedence and Parentheses

To Be Continued:

5.6 Syntax Specification

5.6.1 Grammar

The following provides the Assertion parser production rules (yacc specification) as of revision 46 of the Eiffel reference implementation repository (http://svn.openehr.org/ref_impl_eiffel). Note that because of interdependencies with ADL path and cADL production rules, practical implementations may have to include all production rules in one parser.

```
assertions:
  assertion
| assertions assertion

assertion:
  any_identifier : boolean_expression
| boolean_expression
| any_identifier : error

boolean_expression:
  boolean_leaf
| boolean_node

boolean_node:
  SYM_EXISTS absolute_path
| SYM_EXISTS error
| V_ATTRIBUTE_IDENTIFIER SYM_MATCHES SYM_START_CBLOCK c_primitive
SYM_END_CBLOCK
| SYM_NOT boolean_leaf
| arithmetic_expression = arithmetic_expression
| arithmetic_expression SYM_NE arithmetic_expression
| arithmetic_expression SYM_LT arithmetic_expression
| arithmetic_expression SYM_GT arithmetic_expression
| arithmetic_expression SYM_LE arithmetic_expression
| arithmetic_expression SYM_GE arithmetic_expression
| boolean_expression SYM_AND boolean_expression
| boolean_expression SYM_OR boolean_expression
| boolean_expression SYM_XOR boolean_expression
| boolean_expression SYM_IMPLIES boolean_expression
```

```
boolean_leaf:
  ( boolean_expression )
| SYM_TRUE
| SYM_FALSE

arithmetic_expression:
  arithmetic_leaf
| arithmetic_node

arithmetic_node:
  arithmetic_expression + arithmetic_leaf
| arithmetic_expression - arithmetic_leaf
| arithmetic_expression * arithmetic_leaf
| arithmetic_expression / arithmetic_leaf
| arithmetic_expression ^ arithmetic_leaf

arithmetic_leaf:
  ( arithmetic_expression )
| integer_value
| real_value
| absolute_path
```

5.6.2   Symbols

See the lexical specification for the cADL grammar.

6   Declarations

This section for a future release of ADL

To Be Determined: main problem of variables is that they must have names, which are language-dependent; imagine if there were a mixture of variables added by authors in different languages. The only solution is to name them with terms.

To Be Determined: Variables have to be treated as term coordinations, and should be coded e.g. using ccNNNN codes ("cc" = coordinated code). Then they can be given meanings in any language.

6.1   Predefined Variables

A number of predefined variables can be referenced in ADL assertion expressions, without prior definition, including:

- `$current_date`: `Date`; returns the date whenever the archetype is evaluated
- `$current_time`: `Time`; returns time whenever the archetype is evaluated
- `$current_date_time`: `Date_Time`; returns date/time whenever the archetype is evaluated

To Be Continued: these should be coded as well, using openEHR codes

6.2   Archetype-defined Variables

Variables can also be defined inside an archetype, as part of the assertion statements in an invariant. The syntax of variable definition is as follows:

```
let $var_name = reference
```

Here, a reference can be any of the operand types listed above. 'Let' statements can come anywhere in an invariant block, but for readability, should generally come first.

The following example illustrates the use of variables in an invariant block:

```
invariant
    let $sys_bp =
        /[at0000]/data[at9001]/events[at9002]/data[at1000]/items[at1100]
    let $dia_bp =
        /[at0000]/data[at9001]/events[at9002]/data[at1000]/items[at1200]
    $sys_bp >= $dia_bp
```
To Be Continued:

7 ADL Paths

7.1 Overview

The notion of paths is integral to ADL, and a common path syntax is used to reference nodes in both dADL and cADL sections of an archetype. The same path syntax works for both, because both dADL and cADL have an alternating object/attribute structure. However, the interpretation of path expressions in dADL and cADL differs slightly; the differences are explained in the dADL and cADL sections of this document. This section describes only the common syntax and semantics.

The general form of the path syntax is as follows (see syntax section below for full specification):

```
path: ['/'] { path_segment }+
path_segment: attr_name [ '[' object_id ']' ]
```

Essentially ADL paths consist of segments separated by slashes ('/'), where each segment is an attribute name with optional object identifier, indicated by brackets ('[]').

> **ADL Paths** are formed from an alternation of segments made up of an attribute name and optional object node identifier, separated by slash ('/') characters. Node identifiers are delimited by brackets (i.e. []).

Similarly to paths used in file systems, ADL paths are either absolute or relative, with the former being indicated by a leading slash.

> Paths are **absolute** or **relative** with respect to the document in which they are mentioned. Absolute paths commence with an initial slash ('/') character.

The ADL path syntax also supports the concept of "movable" path patterns, i.e. paths that can be used to find a section anywhere in a hierarchy that matches the path pattern. Path patterns are indicated with a leading double slash ("//") as in Xpath.

> Path **patterns** are absolute or relative with respect to the document in which they are mentioned. Absolute paths commence with an initial slash ('/') character.

7.2 Relationship with W3C Xpath

The ADL path syntax is semantically a subset of the Xpath query language, with a few syntactic shortcuts to reduce the verbosity of the most common cases. Xpath differentiates between "children" and "attributes" sub-items of an object due to the difference in XML between Elements (true subobjects) and Attributes (tag-embedded primitive values). In ADL, as with any pure object formalism, there is no such distinction, and all subparts of any object are referenced in the manner of Xpath children; in particular, in the Xpath abbreviated syntax, the key `child::` does not need to be used.

ADL does not distinguish attributes from children, and also assumes the `node_id` attribute. Thus, the following expressions are legal for cADL structures:

```
items[1]        -- the first member of 'items'
items[systolic] -- the member of 'items' with meaning 'systolic'
items[at0001]   -- the member of 'items' with node id 'at0001'
```

The Xpath equivalents are:

```
items[1]        -- the first member of 'items'
items[meaning() = "systolic"]-- the member of 'items' for which the meaning()
function evaluates to "systolic"
items[@node_id = 'at0001']-- the member of 'items' with key 'at0001'
```

In the above, `meaning()` is a notional function is defined for Xpath in *open*EHR, which returns the rubric for the `node_id` of the current node. Such paths are only for display purposes, and paths used for computing always use the 'at' codes, e.g. `items[at0001]`, for which the Xpath equivalent is `items[@node_id = 'at0001']`.

The ADL movable path pattern is a direct analogue of the Xpath syntax abbreviation for the 'descendant' axis.

7.3   Path Syntax

The ADL path grammar is available as an <u>HTML document</u>. This grammar is implemented and tested using lex (.l file) and yacc (.y file) specifications for in the Eiffel programming environment. The current release of these files is available at <u>http://my.openehr.org/wsvn/ref_impl_eiffel/libraries/common_libs/src/structures/object_graph/path/?rev=0&sc=0</u>. The .l and .y files can easily be converted for use in another yacc/lex-based programming environment.

7.3.1   Grammar

The following provides the ADL path parser production rules (yacc specification) as of revision 46 of the Eiffel reference implementation repository (<u>http://svn.openehr.org/ref_impl_eiffel</u>).

```
input:
  movable_path
| absolute_path
| relative_path
| error

movable_path:
  SYM_MOVABLE_LEADER relative_path

absolute_path:
  / relative_path
| absolute_path / relative_path

relative_path:
  path_segment
| relative_path / path_segment

path_segment:
  V_ATTRIBUTE_IDENTIFIER V_LOCAL_TERM_CODE_REF
| V_ATTRIBUTE_IDENTIFIER
```

7.3.2   Symbols

The following specifies the symbols and lexical patterns used in the path grammar.

```
----------/* symbols */ --------------------------------------------
"."        Dot_code
"/"        Slash_code

"["        Left_bracket_code
"]"        Right_bracket_code

"//"       SYM_MOVABLE_LEADER
```

```
----------/* term code reference */ -----------------------------------
\[[a-zA-Z0-9][a-zA-Z0-9._\-]*\]          V_LOCAL_TERM_CODE_REF

----------/* identifiers */ -------------------------------------------
[A-Z][a-zA-Z0-9_]*                       V_TYPE_IDENTIFIER

[a-z][a-zA-Z0-9_]*[ ]*\(\)               V_FEATURE_CALL_IDENTIFIER

[a-z][a-zA-Z0-9_]*                       V_ATTRIBUTE_IDENTIFIER
```

8 ADL - Archetype Definition Language

This section describes ADL archetypes as a whole, adding a small amount of detail to the descriptions of dADL and cADL already given. The important topic of the relationship of the cADL-encoded `definition` section and the dADL-encoded `ontology` section is discussed in detail. In this section, only standard ADL (i.e. the cADL and dADL constructs and types described so far) is assumed. Archetypes for use in particular domains can also be built with more efficient syntax and domain-specific types, as described in Customising ADL on page 99, and the succeeding sections.

An ADL archetype is a dADL document with the structure shown below:

```
archetype_id = <"some.archetype.id">
adl_version = <"2.0">
is_controlled = <True>
parent_archetype_id = <"some.other.archetype.id">
concept = <[concept_code]>
original_language = <"lang">
translations = <
    ...
>
description = <
    ...
>
definition = (cadl) <#
    cADL plug-in section
#>
invariant = (aadl) <#
    assertions plug-in section
#>
ontology = <
    ...
>
revision_history = <
    ...
>
```

In the above, all top-level attribute names are an exact match in name and type for a corresponding attribute in the *open*EHR. Archetype Object Model (AOM). Optional parts are shown unbolded. In this document, top level attributes are usually described as 'sections' of the archetype.

8.1 Basics

8.1.1 Order of Sections

As of ADL 2.0, order is no longer significant in archetypes. However, it is strongly recommended that the order above be respected, since it is the logical reading order, and some parser implementations may rely on it.

8.1.2 Keywords

ADL has no keywords of its own (i.e, distinct from the keywords in cADL and the invariant language), although depending on how parsers are built, the top-level attribute names might be regarded as special in some implementations.

8.1.3 Node Identification

In the `definition` section of an ADL archetype, a particular scheme of codes is used for node identifiers as well as for denoting constraints on textual (i.e. language dependent) items. Codes are either local to the archetype, or from an external lexicon. This means that the archetype description is the same in all languages, and is available in any language that the codes have been translated to. All term codes are shown in brackets (`[]`). Codes used as node identifiers and defined within the same archetype are prefixed with "at" and by convention have 4 digits, e.g. `[at0010]`. Codes of any length are acceptable in ADL archetypes. Specialisations of locally coded concepts have the same root, followed by "dot" extensions, e.g. `[at0010.2]`. From a terminology point of view, these codes have the implied semantics of subsumption; additionally - the "dot" structuring is used as an optimisation on node identification.

8.1.4 Local Constraint Codes

A second kind of local code is used to stand for constraints on textual items in the body of the archetype. Although these could be included in the main archetype body, because they are language- and/or terminology-sensitive, they are defined in the ontology section, and referenced by codes prefixed by "ac", e.g. `[ac0009]`. As for "at" codes, the convention used in this document is to use 4-digit "ac" codes, even though any number of digits is acceptable. The use of these codes is described in section 8.5.4

8.2 Header Sections

8.2.1 Archetype_id Section

This section introduces the archetype and must include an identifier. A typical `archetype` section is as follows:

```
archetype_id = <"mayo.openehr-ehr-entry.haematology.v1">
```

The multi-axial identifier identifies archetypes in a global space. The syntax of the identifier is described under Archetype Identification on page 9 in The openEHR Archetype System.

8.2.2 Adl_version Section

This mandatory section indicates the version of ADL being used in the archetype. Its value is one of the revision values of this specification, such as "2.0"; it typically looks as follows:

```
adl_version = <"2.0">
```

8.2.3 Is_controlled Section

This is an optional section indicating whether the archetype is change-controlled or not can be included after the version, as follows:

```
is_controlled = <False>
```

The flag may have the values `True` and `False` only, and is an aid to software. Archetypes having `is_controlled` set to True must have the `revision_history` section included, while those with the value `False`, or no flag at all, may omit the `revision_history` section. This enables archetypes to be privately edited in an early development phase without generating large revision histories of little or no value.

8.2.4   Parent_archetype_id Section

This optional section indicates that the archetype is a specialisation of some other archetype, whose identity must be given. Only one specialisation parent is allowed, i.e. an archetype cannot 'multiply-inherit' from other archetypes. An example of declaring specialisation is as follows:

```
archetype_id = <"mayo.openehr-ehr-entry.haematology-cbc.v1">
parent_archetype_id = <"mayo.openehr-ehr-entry.haematology.v1">
```

Here the identifier of the new archetype is derived from that of the parent by adding a new section to its domain concept section. See Archetype Identification on page 9 in The openEHR Archetype System.

8.2.5   Concept Section

All archetypes represent some real world concept, such as a "patient", a "blood pressure", or an "antenatal examination". The concept section describes the overall concept using the code of a concept whose definition is given in the archetype ontology. Like any coded entity, it can be displayed in any language the archetype has been translated to. A typical `concept` section is as follows:

```
concept = <[at0010]>      -- haematology result
```

In this concept definition, the term definition of `[at0010]` contains a more complete description corresponding to the notion implied by the short identifier "`haematology-cbc`" in the example `archeytpe_id` section shown above.

8.2.6   Original_language and Translations Sections

The `original_language` and `translations` sections include meta-data describing the original language in which the archetype was authored (essential for evaluating natural language quality), and the total list of languages available in the archetype due to translation. There can be only one `original_language`. The `translations` section is optional, but must be present if any translations are present in the `description` or `ontology` sections; if present, it must be updated every time a translation of the archetype is undertaken. The following shows a typical example.

```
original_language = <
    <"en">
>
translations = <
    ["de"] = <
        author = <"Frederick Smith">
        accreditation = <"British Medical Translator id 00400595">
    >
    ["ru"] = <
        author = <"Vladimir Korotkov">
        accreditation = <"Russion Translator id 892230A">
        other_details = <
            ["email"] = <"vladimir.korotkov@acme.translators.ru">
        >
    >
>
```

Archetypes must always be translated completely, or not at all, to be valid. This means that when a new translation is made, every language dependent section of the `description` and `ontology` sections has to be translated into the new language, and an appropriate addition made to the `translations` list in the language section.

8.2.7   Description Section

The description section of an archetype contains descriptive information, or what some people think of as document "meta-data", i.e. items that can be used in repository indexes and for searching. The dADL syntax is used for the description, as in the following example.

```
description = <
    original_author = <
        ["name"] = <"Dr J Joyce">
        ["organisation"] = <"NT Health Service">
        ["date"] = <2003-08-03>
    >
    lifecycle_state = <"initial">
    archetype_package_uri =
        <"http://www.aihw.org.au/data_sets/diabetic_archetypes.html">
    details = <
        ["en"] = <
            purpose = <"archetype for diabetic patient review">
            use = <"used for all hospital or clinic-based diabetic reviews,
                including first time. Optional sections are removed according
                to the particular review">
            misuse = <"not appropriate for pre-diagnosis use">
            original_resource_uri =
                <"http://www.healthdata.org.au/data_sets/
                        diabetic_review_data_set_1.html">
            other_details = <...>
        >
        ["de"] = <
            purpose = <"Archetyp für die Untersuchung von Patienten
                    mit Diabetes">
            use = <"wird benutzt für alle Diabetes-Untersuchungen im
                    Krankenhaus, inklusive der ersten Vorstellung. Optionale
                    Abschnitte werden in Abhängigkeit von der speziellen
                    Vorstellung entfernt.">
            misuse = <"nicht geeignet für Benutzung vor Diagnosestellung">
            original_resource_uri =
                <"http://www.healthdata.org.au/data_sets/
                        diabetic_review_data_set_1.html">
            other_details = <...>
        >
    >
>
```

A number of details are worth noting here. Firstly, the free hierarchical structuring capability of dADL is exploited for expressing the "deep" structure of the details section and its subsections. The design of the objects in this section is specified in the AOM. Secondly, the dADL qualified list form is used to allow multiple translations of the purpose and use to be shown. Lastly, empty items such as misuse (structured if there is data) are shown with just one level of empty brackets. The above example shows meta-data based on the CEN MetaKnow standard, with inclusions from the HL7 Templates Proposal [5] and the meta-data of the SynEx and Australian GeHR archetypes.

Which descriptive items are required will depend on the semantic standards imposed on archetypes by health standards organisations and/or the design of archetype repositories and is not specified by ADL.

8.3   Definition Section

The `definition` section contains the main formal definition of the archetype, and is written in the Constraint Definition Language (cADL). A typical `definition` section is as follows:

```
definition = (cadl) <#
    ENTRY[at0000] ∈ {              -- blood pressure measurement
        name ∈ {                   -- any synonym of BP
            CODED_TEXT ∈ {
                code ∈ {
                    CODE_PHRASE ∈ {[ac0001]}
                }
            }
        }
        data ∈ {
            HISTORY[at9001] ∈ {                        -- history
                events cardinality ∈ {1..*} ∈ {
                    EVENT[at9002] occurrences ∈ {0..1} ∈ {-- baseline
                        name ∈ {
                            CODED_TEXT ∈ {
                                code ∈ {
                                    CODE_PHRASE ∈ {[ac0002]}
                                }
                            }
                        }
                        data ∈ {
                            ITEM_LIST[at1000] ∈ {   -- systemic arterial BP
                                items cardinality ∈ {2..*} ∈ {
                                    ELEMENT[at1100] ∈ {          -- systolic BP
                                        name ∈ {       -- any synonym of 'systolic'
                                            CODED_TEXT ∈ {
                                                code ∈ {
                                                    CODE_PHRASE ∈ {[ac0002]}
                                                }
                                            }
                                        }
                                        value ∈ {
                                            QUANTITY ∈ {
                                                magnitude ∈ {|0..1000|}
                                                property ∈ {[properties::0944]}
                                                                    -- "pressure"
                                                units ∈ {[units::387]}  -- "mm[Hg]"
                                            }
                                        }
                                    }
                                    ELEMENT[at1200] ∈ {        -- diastolic BP
                                        name ∈ {       -- any synonym of 'diastolic'
                                            CODED_TEXT ∈ {
                                                code ∈ {
                                                    CODE_PHRASE ∈ {[ac0003]}
                                                }
                                            }
                                        }
                                        value ∈ {
                                            QUANTITY ∈ {
                                                magnitude ∈ {|0..1000|}
                                                property ∈ {[properties::0944]}
                                                                    -- "pressure"
```

```
                                    units ∈ {[units::387]}  -- "mm[Hg]"
                                }
                            }
                        }
                        ELEMENT[at9000] occurrences ∈ {0..*} ∈ {*}
                                                        -- unknown new item
                    }
            ...
```

This definition expresses constraints on instances of the types ENTRY, HISTORY, EVENT, LIST_S, ELE-
MENT, QUANTITY, and CODED_TEXT so as to allow them to represent a blood pressure measurement,
consisting of a history of measurement events, each consisting of at least systolic and diastolic pres-
sures, as well as any number of other items (expressed by the [at9000] "any" node near the bottom).

8.3.1   Design-time and Run-time paths

All non-unique sibling nodes in a cADL text which correspond to nodes in data which might be
referred to from elsewhere in the archetype, or might be used for querying at runtime, require a node
identifier, and it is usually preferable to assign a design-time meaning, enabling paths and queries to
be expressed using logical meanings rather than meaningless identifiers. When data are created
according to a cADL specification, the archetype node identifiers are written into the data, providing
a reliable way of finding data nodes, regardless of what other runtime names might have been chosen
by the user for the node in question. There are two reasons for doing this. Firstly, querying cannot rely
on runtime names of nodes (e.g. names like "sys BP", "systolic bp", "sys blood press." entered by a
doctor are unreliable for querying); secondly, it allows runtime data retrieved from a persistence
mechanism to be re-associated with the cADL structure which was used to create it.

An example which clearly shows the difference between design-time meanings associated with node
ids and runtime names is the following, for the root node of a SECTION headings hierarchy represent-
ing the problem/SOAP headings (a simple heading structure commonly used by clinicians to record
patient contacts under top-level headings corresponding to the patient's problem(s), and under each
problem heading, the headings "subjective", "objective", "assessment", and "plan").

```
SECTION[at0000] matches {                    -- problem
    name matches {
        CODED_TEXT matches {
            code matches {[ac0001]}
        }
    }
}
```

In the above, the node identifier [at0000] is assigned a meaning such as "clinical problem" in the
archetype ontology section. The subsequent lines express a constraint on the runtime *name* attribute,
using the internal code [ac0001]. The constraint [ac0001] is also defined in the archetype ontology
section with a formal statement meaning "any clinical problem type", which could clearly evaluate to
thousands of possible values, such as "diabetes", "arthritis" and so on. As a result, in the runtime data,
the node identifier corresponding to "clinical problem" and the actual problem type chosen at runtime
by a user, e.g. "diabetes", can both be found. This enables querying to find all nodes with meaning
"problem", or all nodes describing the problem "diabetes". Internal [acNNNN] codes are described in
Local Constraint Codes on page 86.

8.4   Invariant Section

The invariant section in an ADL archetype introduces assertions which relate to the entire arche-
type, and can be used to make statements which are not possible within the block structure of the
definition section. Any constraint which relates more than one property to another is in this cate-

gory, as are most constraints containing mathematical or logical formulae. Invariants are expressed in the archetype assertion language, described in section 5 on page 75.

An invariant statement is a first order predicate logic statement which can be evaluated to a boolean result at runtime. Objects and properties are referred to using paths.

The following simple example says that the speed in kilometres of some node is related to the speed-in-miles by a factor of 1.6:

```
invariant = (aadl) <#
    validity: /speed[at0002]/kilometres/magnitude =
              /speed[at0004]/miles/magnitude * 1.6
#>
```

Note that in a well-designed archetype, the '1.6' above should be coded and included in the ontology section.

8.5 Ontology Section

8.5.1 Overview

The `ontology` section of an archetype is expressed in dADL, and is where codes representing node IDs, constraints on text or terms, and bindings to terminologies are defined. Linguistic language translations are added in the form of extra blocks keyed by the relevant language. The following example shows the general layout of this section.

```
ontology = <
    terminologies_available = <"snomed_ct", ...>
    term_definitions = <
        ["en"] = <
            ["at0000"] = <...>
            ...
        >
        ["de"] = <
            ["at0000"] = <...>
            ...
        >
    >
    term_binding = <
        ["snomed_ct"] = <
            ["at0000"] = <...>
            ...
        >
        ...
    >
    constraint_definitions = <
        ["en"] = <...>
            ["ac0001"] = <...>
            ...
        ["de"] = <
            ["ac0001"] = <...>
            ...
        >
        ...
    >
    constraint_binding = <
        ["snomed_ct"] = <...>
```

```
            ...
        >
    >
```

The `term_definitions` section is mandatory, and must be defined for each translation carried out.

Each of these sections can have its own meta-data, which appears within description sub-sections, such as the one shown above providing translation details.

8.5.2   Terminologies_available Sub-section

This section provides the total list of external terminologies which are mentioned anywhere else in the ontology.

```
        terminologies_available = <"snomed_ct", ...>
```

8.5.3   Term_definitions Sub-section

This section is where all archetype local terms (that is, terms of the form [atNNNN]) are defined. The following example shows an extract from the English and German term definitions for the archetype local terms in a problem/SOAP headings archetype. Each term is defined using a structure of name/value pairs, and mustat least include the names "text" and "description", which are akin to the usual rubric, and full definition found in terminologies like SNOMED-CT. Each term object is then included in the appropriate language list of term definitions, as shown in the example below.

```
    term_definitions = <
        ["en"] = <
            ["at0000"] = <
                text = <"problem">
                description = <"The problem experienced by the subject
                    of care to which the contained information relates">
            >
            ["at0001"] = <
                text = <"problem/SOAP headings">
                description = <"SOAP heading structure for multiple problems">
            >
            ...
            ["at4000"] = <
                text = <"plan">
                description = <"The clinician's professional advice">
            >
        >
    >
    ["de"] = <
        ["at0000"] = <
            text = <"klinisches Problem">
            description = <"Das Problem des Patienten worauf sich diese \
                    Informationen beziehen">
        >
        ["at0001"] = <
            text = <"Problem/SOAP Schema">
            description = <"SOAP-Schlagwort-Gruppierungsschema fuer
                    mehrfache Probleme">
        >
        ["at4000"] = <
            text = <"Plan">
            description = <"Klinisch-professionelle Beratung des
                    Pflegenden">
        >
```

>

In some cases, term definitions may have been lifted from existing terminologies. This is only a safe thing to do if the following conditions apply:

- if there is a lexical match between the intended term in the archetype and the preferred term or a synonym in the terminology;
- the definitions *exactly* match the need in the archetype;
- if the classification of the term in the terminology would classify real world instances (e.g. people with a certain illness) the same way as intended by the archetype design.

To indicate where definitions come from, a "provenance" tag can be used, as follows:

```
["at4000"] = <
    text = <"plan">;
    description = <"The clinician's professional advice">;
    provenance = <"ACME_terminology(v3.9a)">
>
```

Note that this does not indicate a *binding* to any term, only its origin. Bindings are described in section 8.5.5 and section 8.5.6.

8.5.4   Constraint_definitions Sub-section

The `constraint_definition` section is of exactly the same form as the `term_definition` section, and provides the definitions - i.e. the meanings - of the local constraint codes, which are of the form `[acNNNN]`. Each such code refers to some constraint such as "any term which is a subtype of 'hepatitis' in the ICD9AM terminology"; the constraint definitions do not provide the constraints themselves, but define the *meanings* of such constraints, in a manner comprehensible to human beings, and usable in GUI applications. This may seem a superfluous thing to do, but in fact it is quite important. Firstly, term constraints can only be expressed with respect to particular terminologies - a constraint for "kind of hepatitis" would be expressed in different ways for each terminology which the archetype is bound to. For this reason, the actual constraints are defined in the `constraint_binding` section. An example of a constraint term definition for the hepatitis constraint is as follows:

```
["at1015"] = <
    text = <"type of hepatitis">
    description = <"any term which means a kind of viral hepatitis">
>
```

Note that while it often seems tempting to use classification codes, e.g. from the ICD vocabularies, these will rarely be much use in terminology or constraint definitions, because it is nearly always *descriptive*, not classificatory terms which are needed.

8.5.5   Term_binding Sub-section

This section is used to describe the equivalences between archetype local terms and terms found in external terminologies. The purpose is solely for allowing query engine software which wants to search for an instance of some external term to determine what equivalent to use in the archetype. Note that this is distinct from the process of embedding mapped terms in runtime data, which is also possible with the data models of HL7v3, openEHR, and CEN 13606.

A typical term binding section resembles the following:

```
term_binding = <
    ["umls"] = <
        ["at0000"] = <[umls::C124305]> -- apgar result
        ["at0002"] = <[umls::0000000]> -- 1-minute event
```

```
        ["at0004"] = <[umls::C234305]> -- cardiac score
        ["at0005"] = <[umls::C232405]> -- respiratory score
        ["at0006"] = <[umls::C254305]> -- muscle tone score
        ["at0007"] = <[umls::C987305]> -- reflex response score
        ["at0008"] = <[umls::C189305]> -- color score
        ["at0009"] = <[umls::C187305]> -- apgar score
        ["at0010"] = <[umls::C325305]> -- 2-minute apgar
        ["at0011"] = <[umls::C725354]> -- 5-minute apgar
        ["at0012"] = <[umls::C224305]> -- 10-minute apgar
    >
>
```

Each entry simply indicates which term in an external terminology is equivalent to the archetype internal codes. Note that not all internal codes necessarily have equivalents: for this reason, a terminology binding is assumed to be valid even if it does not contain all of the internal codes.

8.5.6   Constraint_binding Sub-section

The last of the ontology sections formally describes text constraints from the main archetype body. They are described separately because they are terminology dependent, and because there may be more than one for a given logical constraint. A typical example follows:

```
constraint_binding = <
    ["snomed_ct"] = <
        ["ac0001"] = <http://terminology.org?terminology_id=snomed_ct&&
            has_relation=[102002];with_target=[128004]>
        ["ac0002"] = <http://terminology.org?terminology_id=snomed_ct&&
            synonym_of=[128025]>
    >
>
```

In this example, each local constraint code is formally defined to refer to the result of a query to a service, in this case, a terminology service which can interrogate the Snomed-CT terminology.

8.6   Revision_history Section

The revision_history section of an archetype shows the audit history of changes to the archetype, and is expressed in dADL syntax. It is optional, and is included at the end of the archetype, since it does not contain content of direct interest to archetype authors, and will monotonically grow in size. Where archetypes are stored in a version-controlled repository such as CVS or an equivalent commercial product, the revision history section would normally be regenerated each time by the authoring software, e.g. via processing of the output of the 'prs' command used with SCCS files, or 'rlog' for RCS files. The following shows a typical example, with entries in most-recent-first order (although technically speaking, the order is irrelevant to ADL).

```
revision_history = <
    ["1.57"] = <
        committer = <"Miriam Hanoosh">
        committer_organisation = <"AIHW.org.au">
        time_committed = <2004-11-02 09:31:04+1000>
        revision = <"1.2">
        reason = <"Added social history section">
        change_type = <"Modification">
    >
    -- etc
    ["1.1"] = <
        committer = <"Enrico Barrios">
        committer_organisation = <"AIHW.org.au">
```

```
        time_committed = <2004-09-24 11:57:00+1000>
        revision = <"1.1">
        reason = <"Updated HbA1C test result reference">
        change_type = <"Modification">
    >
    ["1.0"] = <
        committer = <"Enrico Barrios">
        committer_organisation = <"AIHW.org.au">
        time_committed = <2004-09-14 16:05:00+1000>
        revision = <"1.0">
        reason = <"Initial Writing">
        change_type = <"Creation">
    >
>
```

8.7 Validity Rules

This section describes the formal (i.e. checkable) semantics of ADL archetypes. It is recommended that parsing tools use the identifiers published here in their error messages, as an aid to archetype designers.

8.7.1 Global Archetype Validity

The following validity constraints apply to an archetype as a whole. Note that the term "section" means the same as "attribute" in the following, i.e. a section called "definition" in a dADL text is a serialisation of the value for the attribute of the same name.

VARID: archetype identifier validity. The archetype must have an identifier value for the archetype_id section. The identifier must conform to the published *open*EHR specification for archetype identifiers.

VARCN: archetype concept validity. The archetype must have an archetype term value in the concept section. The term must exist in the archetype ontology.

VARDF: archetype definition validity. The archetype must have a definition section, expressed as a cADL syntax string, or in an equivalent plug-in syntax.

VARON: archetype ontology validity. The archetype must have an ontology section, expressed as a cADL syntax string, or in an equivalent plug-in syntax.

VARDT: archetype definition typename validity. The topmost typename mentioned in the archetype definition section must match the type mentioned in the type-name slot of the first segment of the archetype id.

8.7.2 Coded Term Validity

All node identifiers ('at' codes) used in the `definition` part of the archetype must be defined in the `term_definitions` part of the ontology.

VATDF: archetype term validity. Each archetype term used as a node identifier the archetype definition must be defined in the term_definitions part of the ontology.

All constraint identifiers ('ac' codes) used in the `definition` part of the archetype must be defined in the `constraint_definitions` part of the ontology.

VACDF: node identifier validity. Each constraint code used in the archetype definition part must be defined in the constraint_definitions part of the ontology.

8.7.3   Definition Section

The following constraints apply to the `definition` section of the archetype.

> **VDFAI: archetype identifier validity in definition**. Any archetype identifier mentioned in an archetype slot in the definition section must conform to the published *open*EHR specification for archetype identifiers.

> **VDFPT**: **path validity in definition**. Any path mentioned in the definition section must be valid syntactically, and a valid path with respect to the hierarchical structure of the definition section.

9 The ADL Parsing Process

9.1 Overview

FIGURE 6 illustrates the ADL parsing process. An ADL file is converted by the ADL parser into an ADL parse tree. This tree is an in-memory object structure representation of the semantics of the archetype, in a form corresponding to the *open*EHR Archetype Object Model. This model is then validated by the semantic checker of the ADL parser, which can verify numerous things, such as that term codes referenced in the definition section are defined in the ontology section. It can also validate the classes and attributes mentioned in the archetype against a specification for the relevant information model (e.g. in XMI or some equivalent).
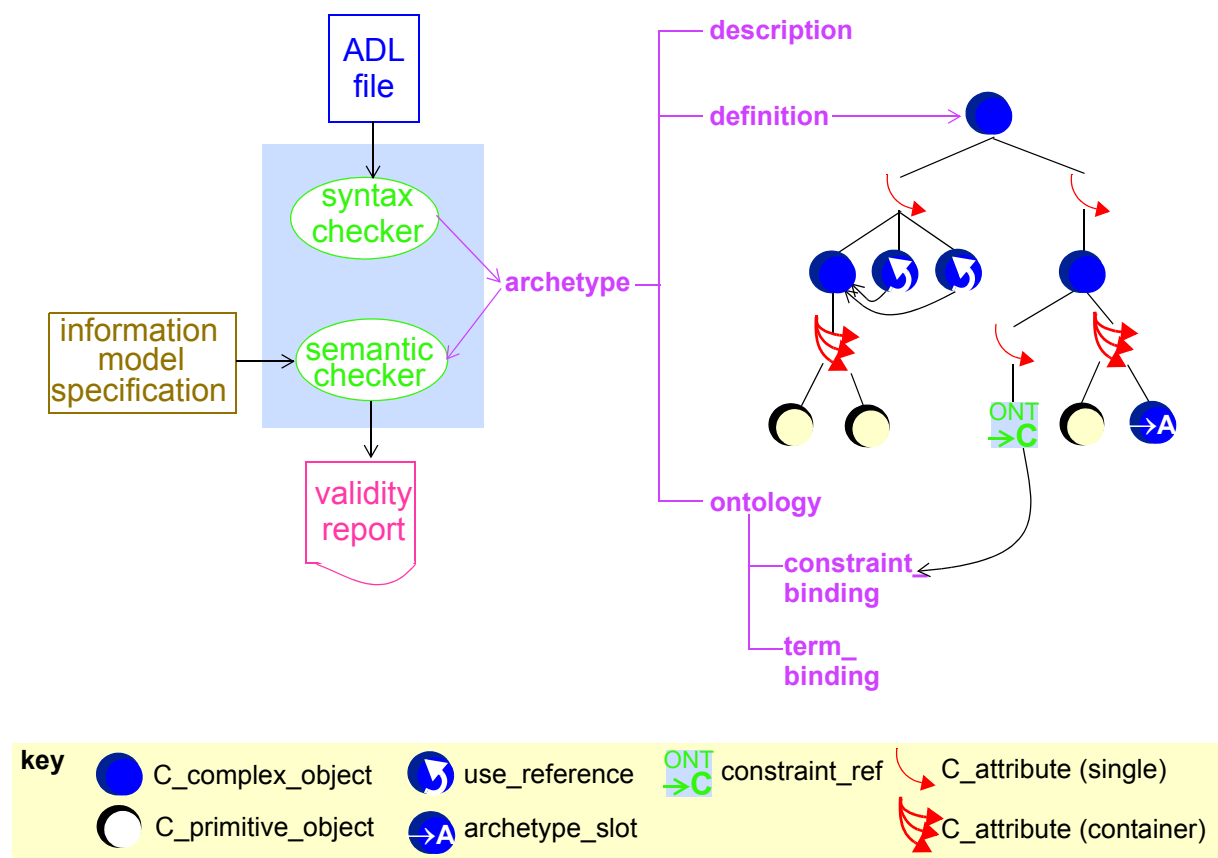

FIGURE 6 Parsed ADL Structure

The object equivalent of an ADL file is shown on the right. It consists of alternate layers of object and relationship nodes, each containing the next level of nodes. At the extremities are leaf nodes - object nodes constraining primitive types such as `String`, `Integer` etc. There are also "use" nodes which represent internal references to other nodes, text constraint nodes which refer to a text constraint in the constraint binding part of the archetype, and archetype constraint nodes, which represent constraints on other archetypes allowed to appear at a given point. The full list of node types is as follows:

C_complex_object: any interior node representing a constraint on instances of some non-primitive type, e.g. `ENTRY`, `SECTION`;

C_attribute: a node representing a constraint on an attribute (i.e. UML 'relationship' or 'primitive attribute') in an object type;

C_primitive_object: a node representing a constraint on a primitive (built-in) object type;

Archetype_internal_ref: a node that refers to a previously defined object node in the same archetype. The reference is made using a path;

Constraint_ref: a node that refers to a constraint on (usually) a text or coded term entity, which appears in the ontology section of the archetype, and in ADL is referred to with an "acNNNN" code. The constraint is expressed in terms of a query on an external entity, usually a terminology or ontology;

Archetype_slot: a node whose statements define a constraint that determines which other archetypes can appear at that point in the current archetype. It can be conceptualised as a keyhole, into which few or many keys might fit, depending on how specific its shape is. Logically it has the same semantics as a C_COMPLEX_OBJECT, except that the constraints are expressed in another archetype, not the current one.

See the *open*EHR Archetype Object Model (AOM) for details.

10   Customising ADL

10.1   Introduction

Standard ADL has a completely regular way of representing constraints. Type names and attribute names from a reference model are mentioned in an alternating, hierarchical structure which is isomorphic to the structure of the corresponding classes in the reference model; constraints at the leaf nodes are represented in a syntactic way which avoids committing to particular modelling details. The overall result enables constraints on most reference model types to be expressed.

10.1.1   Custom Syntax

However, there are occasions for which the standard approach is not enough. One situation is where not everyone in the archetype user base wants to use exactly the same reference model, but nevertheless agrees on the general semantics for many of the types. A typical example of this is the type CODE_PHRASE in the *open*EHR reference model, Data_types package. This type models the notion of a 'coded term', which is ubiquitous in clinical computing. Various user communities in health informatics have slightly different models of the 'coded term' concept, yet all would like to share archetypes which constrain it. This can be achieved by providing additional syntax enabling such constraints to be expressed, while avoiding mentioning any type or attribute names. The following figure shows how this is done, using the example of constraints on the type CODE_PHRASE.

```
                            code matches {
                                CODE_PHRASE matches {
                                    terminology_id matches {"local"}
standard ADL                        code_string matches {"at039"} -- lying
using type and                  }
attribute names                 CODE_PHRASE matches {
                                    terminology_id matches {"local"}
                                    code_string matches {"at040"} -- sitting
                                }
                            }


                            code matches {
clinical ADL                    [local::
syntax                              at039, -- lying
                                    at040] -- sitting
                            }
```

FIGURE 7 Constraints using additional syntax

While these two ADL fragments express exactly the same constraint, the second is clearly shorter and clearer, and avoids implying anything about the formal model of the type of the *code* attribute being constrained.

10.1.2   Custom Constraint Classes

Another situation in which standard ADL falls short is when the required semantics of constraint are different from those provided by the standard approach. Consider a simple type QUANTITY, shown at the top of FIGURE 8, which could be used to represent a person's age in data. A typical ADL constraint to enable QUANTITY to be used to represent age in clinical data is shown below, followed by its

expression in ADL. The only way to do this in ADL is to use multiple alternatives. While this is a per-fectly legal approach, it makes processing by software difficult, since the way such a constraint would be displayed in a GUI would be factored differently.
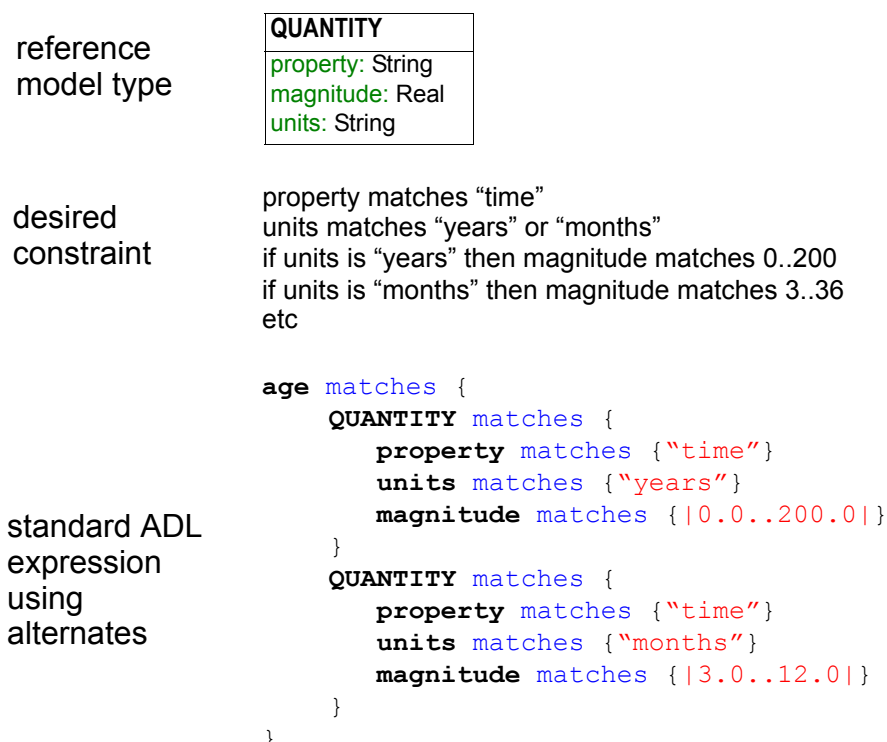
reference model type

```
QUANTITY
property: String
magnitude: Real
units: String
```

desired constraint

property matches "time"
units matches "years" or "months"
if units is "years" then magnitude matches 0..200
if units is "months" then magnitude matches 3..36
etc

standard ADL expression using alternates

```
age matches {
    QUANTITY matches {
        property matches {"time"}
        units matches {"years"}
        magnitude matches {|0.0..200.0|}
    }
    QUANTITY matches {
        property matches {"time"}
        units matches {"months"}
        magnitude matches {|3.0..12.0|}
    }
}
```

FIGURE 8 Standard ADL for Constraint on a Quantity Class

A more powerful possibility is to introduce a new class into the archetype model, representing the concept "constraint on QUANTITY", which we will call C_QUANTITY here. Such a class fits into the class model of archetypes (described in the *open*EHR Archetype Model document), inheriting from the class C_DOMAIN_TYPE. The C_QUANTITY class is illustrated in FIGURE 9, and corresponds to the way constraints on QUANTITY objects are expressed in user applications, which is to say, a property constraint, and a separate list of units/magnitude pairs.
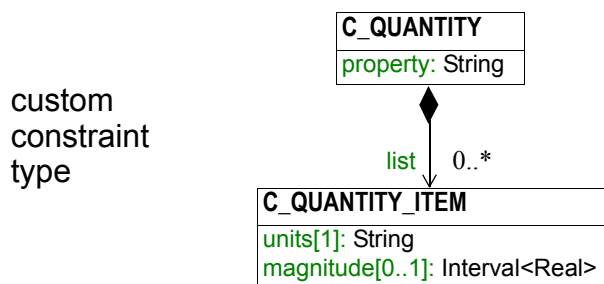
custom constraint type

```
C_QUANTITY
property: String
```

list   0..*

```
C_QUANTITY_ITEM
units[1]: String
magnitude[0..1]: Interval<Real>
```

FIGURE 9 Custom Constraint Type for QUANTITY

The question now is how to express a constraint corresponding to this class in an ADL archetype. The solution is logical, and uses standard ADL. Consider that a particular constraint on a QUANTITY must

be *an instance* of a `C_QUANTITY`; which can be expressed at the appropriate point in the archetype in the form of a section of dADL - the data syntax used in the archetype ontology.

```
value matches {
    C_QUANTITY <
        property = <"time">
        list = <
            items = <
                [1] = <
                    units = <"yr">
                    magnitude = <|0.0..200.0|>
                >
                [2] = <
                    units = <"mth">
                    magnitude = <|1.0..36.0|>
                >
            >
        >
    >
}
```

FIGURE 10 Inclusion of a Constraint Object as Data

This approach can be used for any custom type which represents a constraint on a reference model type. The rules are as follows:

- the dADL section occurs inside the {} block where its standard ADL equivalent would have occurred (i.e. no other delimiters or special marks are needed);
- the dADL section must be 'typed', i.e. it must start with a type name, which should be a rule-based transform of a reference model type (as described in Adding Type Information on page 28);
- the dADL instance must obey the semantics of the custom type of which it is an instance.

It should be understood of course, that just because a custom constraint type has been defined, it does not need to be used to express constraints on the reference model type it targets. Indeed, any mixture of standard ADL and dADL-expressed custom constraints may be used within the one archetype.

11   Relationship of ADL to Other Formalisms

11.1   Overview

Whenever a new formalism is defined, it is reasonable to ask the question: are there not existing formalisms which would do the same job? Research to date has shown that in fact, no other formalism has been designed for the same use, and none easily express ADL's semantics. During ADL's initial development, it was felt that there was great value in analysing the problem space very carefully, and constructing an abstract syntax exactly matched to the solution, rather than attempting to use some other formalism - undoubtedly designed for a different purpose - to try and express the semantics of archetypes, or worse, to start with an XML-based exchange format, which often leads to the conflation of abstract and concrete representational semantics. Instead, the approach used has paid off, in that the resulting syntax is very simple and powerful, and in fact has allowed mappings to other formalisms to be more correctly defined and understood. The following sections compare ADL to other formalisms and show how it is different.

11.2   Constraint Syntaxes

11.2.1   OCL (Object Constraint Language)

The OMG's Object Constraint Language (OCL) appears at first glance to be an obvious contender for writing archetypes. However, its designed use is to write constraints on *object models*, rather than on *data*, which is what archetypes are about. As a concrete example, OCL can be used to make statements about the *actors* attribute of a class `Company` - e.g. that *actors* must exist and contain the `Actor` who is the *lead* of `Company`. However, if used in the normal way to write constraints on a class model, it cannot describe the notion that for a particular kind of (acting) company, such as 'itinerant jugglers', there must be at least four actors, each of whom have among their *capabilities* 'advanced juggling', plus an Actor who has *skill* 'musician'. This is because doing so would constrain all instances of the class `Company` to conform to the specific configuration of instances corresponding to actors and jugglers, when what is intended is to allow a myriad of possibilities. ADL provides the ability to create numerous archetypes, each describing in detail a concrete configuration of instances of type `Company`.

OCL's constraint types include function pre- and post-conditions, and class invariants. There is no structural character to the syntax - all statements are essentially first-order predicate logic statements about elements in models expressed in UML, and are related to parts of a model by 'context' statements. This makes it impossible to use OCL to express an archetype in a structural way which is natural to domain experts. OCL also has some flaws, described by Beale [4].

However, OCL is in fact relevant to ADL. ADL archetypes include invariants (and one day, might include pre- and post-conditions). Currently these are expressed in a syntax very similar to OCL, with minor differences. The exact definition of the ADL invariant syntax in the future will depend somewhat on the progress of OCL through the OMG standards process.

11.3   Ontology Formalisms

11.3.1   OWL (Web Ontology Language)

The Web Ontology Language (OWL) [20] is a W3C initiative for defining Web-enabled ontologies which aim to allow the building of the "Semantic Web". OWL has an abstract syntax [13], developed

at the University of Manchester, UK, and an exchange syntax, which is an extension of the XML-based syntax known as RDF (Resource Description Framework). We discuss OWL only in terms of its abstract syntax, since this is a semantic representation of the language unencumbered by XML or RDF details (there are tools which convert between abstract OWL and various exchange syntaxes).

OWL is a general purpose description logic (DL), and is primarily used to describe "classes" of things in such a way as to support *subsumptive* inferencing within the ontology, and by extension, on data which are instances of ontology classes. There is no general assumption that the data itself were built based on any particular class model - they might be audio-visual objects in an archive, technical documentation for an aircraft or the Web pages of a company. OWL's class definitions are therefore usually constraint statements on an *implied* model on which data *appears* to be based. However, the semantics of an information model can themselves be represented in OWL. Restrictions are the primary way of defining subclasses.

In intention, OWL is aimed at representing some 'reality' and then making inferences about it; for example in a medical ontology, it can infer that a particular patient is at risk of ischemic heart disease due to smoking and high cholesterol, if the knowledge that 'ischemic heart disease has-risk-factor smoking' and 'ischemic heart disease has-risk-factor high cholesterol' are in the ontology, along with a representation of the patient details themselves. OWL's inferencing works by subsumption, which is to say, asserting either that an 'individual' (OWL's equivalent of an object-oriented instance or a type) conforms to a 'class', or that a particular 'class' 'is-a' (subtype of another) 'class'; this approach can also be understood as category-based reasoning or set-containment.

ADL can also be thought of as being aimed at describing a 'reality', and allowing inferences to be made. However, the reality it describes is in terms of constraints on information structures (based on an underlying information model), and the inferencing is between data and the constraints. Some of the differences between ADL and OWL are as follows.

- ADL syntax is predicated on the existence of existing object-oriented reference models, expressed in UML or some similar formalism, and the constraints in an ADL archetype are in relation to types and attributes from such a model. In contrast, OWL is far more general, and requires the explicit expression of a reference model in OWL, before archetype-like constraints can be expressed.

- Because information structures are in general hierarchical compositions of nodes and elements, and may be quite deep, ADL enables constraints to be expressed in a structural, nested way, mimicking the tree-like nature of the data it constrains. OWL does not provide a native way to do this, and although it is possible to express approximately the same constraints in OWL, it is fairly inconvenient, and would probably only be made easy by machine conversion from a visual format more or less like ADL.

- As a natural consequence of dealing with heavily nested structures in a natural way, ADL also provides a path syntax, based on Xpath [21], enabling any node in an archetype to be referenced by a path or path pattern. OWL does not provide an inbuilt path mechanism; Xpath can presumably be used with the RDF representation, although it is not yet clear how meaningful the paths would be with respect to the named categories within an OWL ontology.

- ADL also natively takes care of disengaging natural language and terminology issues from constraint statements by having a separate ontology per archetype, which contains 'bindings' and language-specific translations. OWL has no inbuilt syntax for this, requiring such semantics to be represented from first principles.

- Lastly, OWL (as of mid 2004) is still under development, and has only a very limited set of primitive constraint types (it is not possible for example to state a constraint on an Integer attribute of the form 'any value between 80 and 110'), although this is being addressed; by contrast, ADL provides a rich set of constraints on primitive types, including dates and times.

Research to date shows that the semantics of an archetype are likely to be representable inside OWL, assuming expected changes to improve its primitive constraint types occur. To do so would require the following steps:

- express the relevant reference models in OWL (this has been shown to be possible);
- express the relevant terminologies in OWL (research on this is ongoing);
- be able to represent concepts (i.e. constraints) independently of natural language (status unknown);
- convert the cADL part of an archetype to OWL; assuming the problem of primitive type constraints is solved, research to date shows that this should in principle be possible.

To *use* the archetype on data, the data themselves would have to be converted to OWL, i.e. be expressed as 'individuals'. In conclusion, we can say that mathematical equivalence between OWL and ADL is probably provable. However, it is clear that OWL is far from a convenient formalism to express archetypes, or to use them for modelling or reasoning against data. The ADL approach makes use of existing UML semantics and existing terminologies, and adds a convenient syntax for expressing the required constraints. It also appears fairly clear that even if all of the above conversions were achieved, using OWL-expressed archetypes to validate data (which would require massive amounts of data to be converted to OWL statements) is unlikely to be anywhere near as efficient as doing it with archetypes expressed in ADL or one of its concrete expressions.

Nevertheless, OWL provides a very powerful generic reasoning framework, and offers a great deal of inferencing power of far wider scope than the specific kind of 'reasoning' provided by archetypes. It appears that it could be useful for the following archetype-related purposes:

- providing access to ontological resources while authoring archetypes, including terminologies, pure domain-specific ontologies, etc;
- providing a semantic 'indexing' mechanism allowing archetype authors to find archetypes relating to specific subjects (which might not be mentioned literally within the archetypes);
- providing inferencing on archetypes in order to determine if a given archetype is subsumed within another archetype which it does not specialise (in the ADL sense);
- providing access to archetypes from within a semantic Web environment, such as an ebXML server or similar.

Research on these areas is active in the US, UK, Australia, Spain, Denmark and Turkey(mid 2004).

11.3.2   KIF (Knowledge Interchange Format)

The Knowledge Interchange Format (KIF) is a knowledge representation language whose goal is to be able to describe formal semantics which would be sharable among software entities, such as information systems in an airline and a travel agency. An example of KIF (taken from [10]) used to describe the simple concept of "units" in a QUANTITY class is as follows:

```
(defrelation BASIC-UNIT
    (=> (BASIC-UNIT ?u) ; basic units are distinguished
        (unit-of-measure ?u))) ; units of measure
```

```
(deffunction UNIT*
    ; Unit* maps all pairs of units to units
    (=> (and (unit-of-measure ?u1)
             (unit-of-measure ?u2))
      (and (defined (UNIT* ?u1 ?u2))
             (unit-of-measure (UNIT* ?u1 ?u2))))
; It is commutative
(= (UNIT* ?u1 ?u2) (UNIT* ?u2 ?u1))
; It is associative
(= (UNIT* ?u1 (UNIT* ?u2 ?u3))
    (UNIT* (UNIT* ?u1 ?u2) ?u3)))

(deffunction UNIT^
    ; Unit^ maps all units and reals to units
    (=> (and (unit-of-measure ?u)
          (real-number ?r))
      (and (defined (UNIT^ ?u ?r))
             (unit-of-measure (UNIT^ ?u ?r))))
; It has the algebraic properties of exponentiation
(= (UNIT^ ?u 1) ?u)
(= (unit* (UNIT^ ?u ?r1) (UNIT^ ?u ?r2))
    (UNIT^ ?u (+ ?r1 ?r2)))
(= (UNIT^ (unit* ?u1 ?u2) ?r)
    (unit* (UNIT^ ?u1 ?r) (UNIT^ ?u2 ?r)))
```

It should be clear from the above that KIF is a definitional language - it defines all the concepts it mentions. However, the most common situation in which we find ourselves is that information models already exist, and may even have been deployed as software. Thus, to use KIF for expressing archetypes, the existing information model and relevant terminologies would have to be converted to KIF statements, before archetypes themselves could be expressed. This is essentially the same process as for expressing archetypes in OWL.

It should also be realised that KIF is intended as a knowledge exchange format, rather than a knowledge representation format, which is to say that it can (in theory) represent the semantics of any other knowledge representation language, such as OWL. This distinction today seems fine, since Web-enabled languages like OWL probably don't need an exchange format other than their XML equivalents to be shared. The relationship and relative strengths and deficiencies is explored by e.g. Martin [11].

11.4   XML-based Formalisms

11.4.1   XML-schema

Previously, archetypes have been expressed as XML instance documents conforming to W3C XML schemas, for example in the Good Electronic Health Record (GeHR; see http://www.gehr.org) and *open*EHR projects. The schemas used in those projects correspond technically to the XML expressions of information model-dependent object models shown in FIGURE 2. XML archetypes are accordingly equivalent to serialised instances of the parse tree, i.e. particular ADL archetypes serialised from objects into XML instance.

A   References

Publications

1   Beale T. *Archetypes: Constraint-based Domain Models for Future-proof Information Systems*. OOPSLA 2002 workshop on behavioural semantics.
Available at http://www.deepthought.com.au/it/archetypes.html.

2   Beale T. *Archetypes: Constraint-based Domain Models for Future-proof Information Systems*. 2000. Available at http://www.deepthought.com.au/it/archetypes.html.

3   Beale T, Heard S. The openEHR Archetype Object Model. See http://www.openehr.org/repositories/spec-dev/latest/publishing/architecture/archetypes/archetype_model/REV_HIST.html.

4   Beale T. *A Short Review of OCL*.
See http://www.deepthought.com.au/it/ocl_review.html.

5   Dolin R, Elkin P, Mead C *et al*. *HL7 Templates Proposal*. 2002.
Available at http://www.hl7.org.

6   Heard S, Beale T. Archetype Definitions and Principles. See http://www.openehr.org/repositories/spec-dev/latest/publishing/architecture/archetypes/principles/REV_HIST.html.

7   Heard S, Beale T. *The openEHR Archetype System*. See http://www.openehr.org/repositories/spec-dev/latest/publishing/architecture/archetypes/system/REV_HIST.html.

8   Hein J L. *Discrete Structures, Logic and Computability (2nd Ed)*. Jones and Bartlett 2002.

9   Kilov H, Ross J. *Information Modelling: an Object-Oriented Approach*. Prentice Hall 1994.

10   Gruber T R. *Toward Principles for the Design of Ontologies Used for Knowledge Sharing*. in Formal Ontology in Conceptual Analysis and Knowledge Representation. Eds Guarino N, Poli R. Kluwer Academic Publishers. 1993 (Aug revision).

11   Martin P. Translations between UML, OWL, KIF and the WebKB-2 languages (For-Taxonomy, Frame-CG, Formalized English). May/June 2003. Available at http://meganesia.int.gu.edu.au/~ph-martin/WebKB/doc/model/comparisons.html as at Aug 2004.

12   Meyer B. *Eiffel the Language (2nd Ed)*. Prentice Hall, 1992.

13   Patel-Schneider P, Horrocks I, Hayes P. *OWL Web Ontology Language Semantics and Abstract Syntax*.
See http://w3c.org/TR/owl-semantics/.

14   Smith G. *The Object Z Specification Language*. Kluwer Academic Publishers 2000. See http://www.itee.uq.edu.au/~smith/objectz.html.

15   Sowa J F. *Knowledge Representation: Logical, philosophical and Computational Foundations*. 2000, Brooks/Cole, California.

Resources

16   HL7 v3 RIM. See http://www.hl7.org.

17   *open*EHR. EHR Reference Model. See http://svn.openehr.org/specification/TRUNK/project_page.htm.

18   Perl Regular Expressions. See http://www.perldoc.com/perl5.6/pod/perlre.html.

19   SynEx project, UCL. http://www.chime.ucl.ac.uk/HealthI/SynEx/.

20   W3C. *OWL - the Web Ontology Language*.
See http://www.w3.org/TR/2003/CR-owl-ref-20030818/.

21   W3C. XML Path Language. See http://w3c.org/TR/xpath.

END OF DOCUMENT